

BIMS-LEGAL: Dual-Store Soft Binding for Recovering Prior Legal Advice under Same-Domain Interference

Anonymous Author(s)

Abstract

Recovering prior client-specific advice from a legal consultation archive is difficult because dense retrieval may rank a relevant question turn highly while omitting the lawyer’s answer—an *episode incompleteness* failure under same-domain interference. BIMS-LEGAL is a dual-store system for Chinese legal consultation logs: a fast episodic turn index with session identifiers and a slower BIRCH semantic store. Soft O2 performs turn-level session binding by giving siblings attenuated inherited scores (βs) rather than hard score copy. We evaluate exact replay as diagnostic and use paraphrase, follow-up, and advice-recall as primary non-exact channels. Primary Soft O2 contrasts use a shared turn-level index against dense FlatIP, hard hydration, and shuffled-*sid* controls; Soft O2 is strongest on non-exact channels where a partial episode hit can bind its answer sibling. Soft O2-C and gated Hybrid are reported as exploratory, gold-assisted solvability bounds for cross-session recovery—complementary to Soft O2, not replacements for it, and not claims about natural unsupervised BIRCH deployment. We report corrected graded nDCG alongside Answer Hit and Episode Completeness, with paired significance tests and a failure taxonomy separating complete, incomplete, and missed-session retrieval.

Keywords: legal information retrieval, consultation memory, soft episode binding, dual-store memory, cluster binding, Chinese legal corpora

1. Introduction

2 A junior associate may ask a firm assistant: “Last month you advised
3 this client on terminating a fixed-term contract—what remedy did we rec-

Email address: `anonymous@anonymous.edu` (Anonymous Author(s))

ommend?” The system returns several near-identical “contract termination” questions, ranks the stored user utterance highly, and never surfaces the partner’s written advice. The subsequent reply may sound fluent yet be grounded in the wrong episode—or in none. In consultation archives, questions and answers share specialized terminology; neighbouring matters are topically adjacent yet legally distinct; and the professionally useful evidence unit is the full consultation *episode*.

Legal informatics has long studied statutory interpretation, case retrieval, and advisory systems Martinez-Gil (2023); Ashley (2017). Recent legal large language models (LLMs) improve statutory and fact-pattern answering Yue et al. (2023); Huang et al. (2023); Yue et al. (2024). Classical legal information retrieval (IR) mainly targets *external* authority—statutes, cases, or community answers—as documents Askari et al. (2024); Louis et al. (2024); Liu et al. (2022); Zhang et al. (2024). Memory-augmented dialogue systems recover evidence from evolving conversation logs Wu et al. (2024); Maharana et al. (2024); Packer et al. (2023). The gap addressed here is *internal consultation memory*: retrieving prior client-specific advice from a growing same-domain archive under lexical interference.

BIMS-LEGAL is a dual-store Brain-Inspired Memory System for legal consultation archives, designed and evaluated in this study. The *systemic contribution* is a dual-store architecture that separates a fast episodic pathway (timestamped turn embeddings with speaker role and session identifier *sid*) from a slow semantic pathway based on Balanced Iterative Reducing and Clustering using Hierarchies (BIRCH) with deferred consolidation Zhang et al. (1996). Complementary Learning Systems (CLS) theory McClelland et al. (1995); Kumaran et al. (2016) supplies an architectural analogy for this dual-pathway design, not an empirical claim about biological memory. Under same-domain interference, dense turn retrieval in a vector database often returns a question turn while its answer sibling remains outside top- k —a system-level episodic memory failure we term *episode incompleteness*, distinct from a full session miss and exacerbated by IVFPQ score collapse and context fragmentation at moderate archive sizes.

The *algorithmic contributions* are soft binding operators that complete episodes without hard score copy:

- **Soft O2 (session soft binding).** When a turn is retrieved with score s , siblings sharing *sid* inherit $\max(s_t, \beta \cdot s)$, surfacing answer turns under question-hit / answer-miss while preserving rank competition

41 among candidates.

- 42 • **Soft O2-C and gated Hybrid (cluster soft binding).** The same
43 soft-inheritance rule applies within BIRCH clusters when gold evidence
44 is not co-session with the query; Hybrid triggers cluster binding only
45 on direct dense hits to avoid displacing same-session siblings.

46 Two **implementation and scale design operators** support rigorous large-
47 scale shared-store experiments but are not standalone algorithmic contribu-
48 tions:

- 49 • **O1 (exact FlatIP indexing).** At hundreds-to-thousands of turns,
50 IVFPQ product quantization is under-trained and collapses Answer
51 Hit; FlatIP restores faithful turn scores so Soft O2 inherits uncorrupted
52 sibling signals.
- 53 • **O3 (bulk-load consolidation).** Deferred BIRCH and disk writes
54 make archive construction tractable without changing retrieval seman-
55 tics.

56 We ask four research questions:

- 57 • **RQ1.** Under same-domain interference, does episode incompleteness
58 (question hit / answer miss) dominate full session miss?
- 59 • **RQ2.** Do Soft O2 and Soft O2-C address episode incompleteness under
60 measured index distortion, and what incremental gains does Soft O2
61 provide over FlatIP and hard parent hydration?
- 62 • **RQ3.** On the *same* LegalEp /LegalMem stores, when is the slow
63 (BIRCH) pathway effective? Does Soft O2-C help under standard same-
64 session gold, and under a fair Mix of same-/cross-session gold?
- 65 • **RQ4.** Across CAIL multi-turn and LegalEp paraphrase /advice-recall
66 channels, how large and session-tied are Soft O2 gains, and how do
67 Soft O2 and Soft O2-C /Hybrid complement each other?

68 *Contributions..*

- 69 1. A dual-store memory architecture for Chinese legal consultation archives,
70 separating fast episodic turn recovery from slower BIRCH semantic in-
71 tegration under same-domain interference.

- 72 2. Soft O2, a turn-level session soft-binding operator that addresses episode
73 incompleteness without hard parent hydration or session-max score
74 copy—the primary algorithmic claim before Soft O2-C.
- 75 3. A conservative three-candidate diagnostic of Soft O2 score competition
76 using the gap ratio $\rho = s_d/s_h$; default $\beta=0.98$ is chosen by validation
77 sweep, not by a closed-form optimum.
- 78 4. Turn-level lexical and fusion controls (BM25-turn, dense+BM25 RRF,
79 cross-encoder) under the same evidence unit as Soft O2; joint episode
80 / document-level indexes are a different input granularity and are not
81 claimed as Soft O2 competitors here.
- 82 5. Exploratory Soft O2-C / gated Hybrid results under a gold-assisted
83 Mix solvability bound, plus corrected nDCG, failure taxonomy, and
84 paired significance reporting.
- 85 6. O1 exact FlatIP and O3 bulk-load consolidation as implementation
86 operators enabling the shared-store ablations reported here.

87 Section 2 reviews related work. Section 3 presents BIMS-LEGAL opera-
88 tors and the Soft O2 diagnostic model. Section 4 details corpora, corrected
89 metrics, and protocols. Section 5 reports results. Section 6 discusses impli-
90 cations and limitations. Section 7 concludes.

91 2. Related Work

92 2.1. Legal IR and long-horizon dialogue memory

93 Legal QA and retrieval systems typically ground generation in statutes,
94 cases, or published advice Louis et al. (2024); Askari et al. (2024); Martinez-
95 Gil (2023); Ashley (2017). Conversational legal case retrieval and event-aware
96 similar-case ranking further show how interaction context and structured le-
97 gal knowledge shape IR effectiveness Liu et al. (2022); Zhang et al. (2024).
98 Those pipelines optimize access to external authority. Consultation memory
99 retrieval targets a different evidence source: *prior consultation episodes* inside
100 an agent’s own archive. Confusing two clients’ histories can produce fluent
101 yet professionally inappropriate responses even when statute retrieval is cor-
102 rect. Domain LLMs such as DISC-LawLLM, Lawyer LLaMA, and LawLLM
103 improve statutory and advisory generation Yue et al. (2023); Huang et al.
104 (2023); Yue et al. (2024), but they still need a reliable memory index when
105 prior advice must be recovered from a growing same-domain log.

Legal consultation archives add a further difficulty: high lexical overlap among topically adjacent matters, with the professionally useful evidence unit being the full question–answer episode. This setting motivates *dual-store* memory architectures that keep fast episodic turn retrieval separate from slower semantic clustering, and session-aware soft binding operators that recover answer turns under same-domain interference without collapsing rank competition among unrelated candidates.

2.2. Dual-store memory and episodic retrieval

Dual-store designs distinguish fast episodic encoding from slower semantic integration McClelland et al. (1995); Kumaran et al. (2016); Tulving et al. (1972); Atkinson and Shiffrin (1971). Memory architectures for dialogue agents extend context via hierarchical paging, summary banks, graphs, and neuro-symbolic retrieval Packer et al. (2023); Zhong et al. (2024); Rasmussen et al. (2025); Gutiérrez et al. (2024); Modarressi et al. (2023); Wang et al. (2024); Lewis et al. (2020); Wu et al. (2025), but most treat retrieved units as flat passages rather than *multi-turn consultation episodes* whose answer turn may rank below an unrelated question turn sharing legal vocabulary. LongMemEval and LoCoMo evaluate long-horizon recall and QA over personal conversation logs Wu et al. (2024); Maharana et al. (2024), while broader memory-agent surveys emphasize persistence beyond the context window Wu et al. (2025); Tan et al. (2025). Recent conversational RAG work stresses reuse of historical search evidence across turns Mo et al. (2026). BIMS-LEGAL positions its dual-store contribution at this intersection: episodic turn embeddings with *sid* for same-session recovery (Soft O2), and BIRCH centroids with Soft O2-C for cross-session recovery when query and gold evidence occupy different sessions but share thematic structure Zhang et al. (1996). Soft O2 directly addresses the episodic retrieval challenge that parent-document hydration and session-max aggregation handle via hard score copy: under dense same-domain interference, a question turn can dominate top- k while its answer sibling remains absent, producing fluent yet ungrounded downstream generation unless siblings enter under attenuated inherited scores.

2.3. Session-aware retrieval

Parent-document hydration, session aggregation, and joint session indexing are standard session-aware mechanisms for attaching sibling context to a dense hit. Soft episode binding (Soft O2) inherits $\beta \cdot s$ for siblings and keeps

ranking competition among candidates. Parent-document and hierarchical retrieval similarly attach child evidence to a parent unit Callan et al. (1994); Dai et al. (2019); Zaheer et al. (2017); Karpukhin et al. (2020); Soft O2 differs by using attenuated turn-level score propagation rather than hard hydration or atomic session documents. Lexical BM25 Robertson and Zaragoza (2009), dense passage retrieval Karpukhin et al. (2020), dense-sparse Reciprocal Rank Fusion Cormack et al. (2009), and cross-encoder reranking Nogueira et al. (2019); Chen et al. (2024) provide complementary controls outside session membership. Exact flat indexes and product-quantized ANN backends Johnson et al. (2019) further affect whether sibling binding is applied over faithful turn scores. Table 1 summarizes the session-structure contrasts used in the main grids.

Table 1: Session-aware mechanisms compared in this work.

Mechanism	Unit	Expansion	Type	Rank
Parent hydration	Session	Insert siblings	Hard	No
Session-max	Session	Session max	Hard	Same*
Joint episode index	Session doc	Atomic	N/A	N/A
Soft O2	Turn+ <i>sid</i>	Soft β_s	Soft	Yes
Soft O2-C	Turn+cluster	Soft β_{cs}	Soft	Yes

*On the reported corpora, session-max rankings coincide with hard parent hydration under full-sibling expansion; main tables report a single hard-expansion baseline. Soft O2 denotes session soft binding; Soft O2-C denotes cluster soft binding on the BIRCH pathway.

We reuse CLS vocabulary as an architectural analogy for BIMS-LEGAL: turn embeddings with timestamps, roles, and session identifiers as episodic elements; Soft O2 as episode binding on the episodic pathway; BIRCH centroids with Soft O2-C as semantic integration on the slow pathway; and shared-store distractors as interference. The dual-store split is the systemic design; Soft O2 and Soft O2-C are the algorithmic mechanisms evaluated on Chinese legal consultation archives.

3. Method: BIMS-LEGAL

3.1. Architecture overview

Figure 1 summarizes the pipeline. BIMS-LEGAL maintains complementary episodic and semantic stores over the same consultation archive.




Figure 1: BIMS-LEGAL dual-store pipeline. (A) Episodic turn indexing and BIRCH semantic clustering (O3 bulk-load consolidation). (B) Query-time dense retrieval under O1 FlatIP with Soft O2 session binding and Soft O2-C cluster binding (algorithmic core); O1 and O3 are implementation operators. (C) Evaluation order: shared-store ablations, CAIL /LegalEp Soft O2 grids, same-store cluster validation, and fair Mix Soft O2-C.

169 **Definition 1** (Episodic and semantic elements). *An episodic element is*
 170 $e_i = (\mathbf{v}_i, t_i, c_i, \phi_i, \text{sid}_i)$, *where* $\mathbf{v}_i \in \mathbb{R}^d$ *is an embedding,* t_i *a timestamp,* c_i
 171 *conversational context (including speaker role),* ϕ_i *surface features, and* sid_i
 172 *the consultation session identifier. A semantic element is* $s_j = (\mathbf{c}_j, C_j, \kappa_j, \tau_j)$
 173 *with centroid* $\mathbf{c}_j$, *member set* C_j , *label* κ_j , *and formation statistics* τ_j . Write
 174 $\text{Sib}(t) = \{t' : \text{sid}_{t'} = \text{sid}_t\} \setminus \{t\}$ *for session siblings and* $\text{Clus}(t) = \{t' : \text{cid}_{t'} =$
 175 $\text{cid}_t\} \setminus \{t\}$ *for BIRCH-cluster siblings.*

176 **Definition 2** (Dual-store retrieval map). *Given query* q *with embedding*
 177 $\mathbf{v}_q$, *base episodic retrieval returns a scored candidate set* $\mathcal{R}_0(q) = \{(t, s_t)\}$
 178 *under Eq. (1). Soft O2 produces* $\mathcal{R}_{\text{O2}}(q) = \text{SoftBind}(\mathcal{R}_0(q); \text{Sib}, \beta)$ *(Sec-*
 179 *tion 3.3). Soft O2-C produces* $\mathcal{R}_{\text{O2C}}(q) = \text{SoftBind}(\mathcal{R}_0(q); \text{Clus}, \beta_c)$ *(Sec-*
 180 *tion 3.5). Gated Hybrid composes session binding then cross-session cluster*
 181 *binding (Algorithm 2). The returned list is the top-k of the bound set under*
 182 *the rewritten scores.*

183 3.2. Base scoring and implementation operators (O1, O3)

184 For query q with embedding $\mathbf{v}_q$, episodic base retrieval scores a memory
185 turn m by

$$\text{score}(m) = \alpha S(\mathbf{v}_q, \mathbf{v}_m) + (1 - \alpha) R(t_m), \quad (1)$$

186 where S is cosine similarity on ℓ_2 -normalized embeddings and

$$R(t_m) = \max\left(0, 1 - \frac{\Delta_t}{30 \text{ days}}\right), \quad (2)$$

187 with Δ_t the elapsed time between query time and the turn timestamp. Unless
188 stated otherwise, $\alpha=0.7$. Eq. (1) is the ranking score used throughout the
189 legal grids; Soft O2 /Soft O2-C rewrite sibling scores after this fusion step
190 and do not replace Eq. (1).

191 *O1: Exact FlatIP indexing..* Let $\hat{S}$ denote the similarity induced by an ANN
192 backend. At hundreds to a few thousand turns, FAISS IVFPQ is often under-
193 trained, so $\|\hat{S} - S\|$ is large enough to reorder turns and break Soft O2
194 inheritance. O1 forces $\hat{S} = S$ via exact FlatIP (inner-product search on
195 ℓ_2 -normalized embeddings). Quantization error depresses Answer Hit (Ta-
196 ble A.8: IVFPQ AH 0.330 at $M=1600$ on LegalEp-DISC exact). O1 is an
197 engineering prerequisite for rigorous Soft O2 evaluation, not a standalone
198 retrieval algorithm.

199 *O3: Bulk-load consolidation..* During bulk construction, O3 replaces the per-
200 insert update ($e_i \mapsto \text{BIRCH}(e_i); \text{Flush}$) by a deferred operator $\text{Finalize} =$
201 $\text{BIRCH}(\{e_i\}_{i=1}^N) \circ \text{Flush}$, executed once after all episodic writes. O3 does
202 not change retrieval semantics; it makes $M \geq 400$ shared-store experiments
203 runnable (≈ 11 min wall-clock for $M=400$; amortized ≈ 0.825 s/turn over
204 ~ 800 turns). Quality–latency curves to $M=6400$ appear in Section 5.6.

205 3.3. Soft O2: session soft binding

206 *Motivation..* Under same-domain interference, a retrieved question turn can
207 score highly while its answer sibling remains outside top- k —episode incom-
208 pleteness compounded by IVFPQ score collapse and context fragmentation.
209 Soft O2 completes episodes without hard score copy.

210 **Definition 3** (Soft binding operator). *Given a scored set $\mathcal{R} = \{(t, s_t)\}$, a*
211 *sibling relation $\mathcal{N}(\cdot)$, and attenuation $\beta \in (0, 1]$, define*

$$\text{SoftBind}(\mathcal{R}; \mathcal{N}, \beta) : \quad s_{t'} \leftarrow \max(s_{t'}, \beta \cdot s_t) \quad \text{for all } t \in \mathcal{R}, t' \in \mathcal{N}(t). \quad (3)$$

Algorithm 1 Soft O2 session soft binding

Require: Query q , base hits $\mathcal{R}_0$, top- k , attenuation β

```
1:  $\mathcal{R} \leftarrow \mathcal{R}_0$ 
2: for each hit  $t \in \mathcal{R}_0$  with score  $s_t$  do
3:   for each sibling  $t' \in \text{Sib}(t)$  absent from  $\mathcal{R}$  do
4:     insert  $(t', \beta \cdot s_t)$  into  $\mathcal{R}$ 
5:   end for
6: end for
7:                                      $\triangleright$  Completeness pass after fusion with Eq. (1)
8: for each session  $sid$  represented in  $\mathcal{R}$  do
9:    $s_{\max} \leftarrow \max\{s_u : u \in \mathcal{R}, sid_u = sid\}$ 
10:  for each  $t' \in \text{Sib}$  of  $sid$  still absent do
11:    insert  $(t', \beta \cdot s_{\max})$  into  $\mathcal{R}$ 
12:  end for
13: end for
14: return top- $k$  of  $\mathcal{R}$  by descending score
```

212 *Soft O2 is $\text{SoftBind}(\cdot; \text{Sib}, \beta)$ with default $\beta=0.98$. Hard parent hydration is*
213 *the special case $\beta=1$ with forced insertion of all siblings; session-max replaces*
214 *each member score by $\max_{t \in sid} s_t$.*

215 Algorithm 1 states the deployed Soft O2 procedure (expand + complete-
216 ness truncate), matching the public implementation.

217 3.4. Soft O2 attenuation as diagnostic score competition

218 The attenuation β controls how strongly a retrieved turn promotes its
219 session siblings. We treat the following as a *diagnostic* model, not a proof
220 of a universal optimum. In particular, a zero-margin hinge objective with
221 $\beta \in (0, 1]$ does not identify a meaningful quantile solution: when the rank
222 margin is zero, the rank term is inactive for all $\beta < 1$, so its weight cannot
223 determine a closed-form optimum.

224 **Definition 4** (Three-candidate diagnostic model). *For a query whose dense*
225 *retrieval directly hits a gold turn h with fused score $s_h > 0$, let a be the gold*
226 *answer sibling and let d be the strongest non-session distractor in a wide*
227 *candidate pool, with scores s_a and s_d . Soft O2 rewrites*

$$s'_h = s_h, \quad s'_a = \max(s_a, \beta s_h), \quad s'_d = s_d, \quad (4)$$

Table 2: Diagnostic gap-ratio statistics on LegalEp $M=3000$ stores (exact-channel query proxy; strongest measured non-session distractor). These are calibration diagnostics, not derivations of an optimal β .

Corpus	ρ_{50}	ρ_{95}	Cover@0.90	Default β
LegalEp-DISC	0.746	0.854	0.981	0.98
LegalEp-Lawyer	0.662	0.815	0.998	0.98

and defines the gap ratio $\rho = s_d/s_h$. Because d is the strongest measured distractor, conditions based on ρ are conservative sufficient conditions for beating that distractor, not exact top- k boundary conditions for every competitor.

The binding analysis is most relevant in the soft-injection regime $s_a \leq \beta s_h$, where $s'_a = \beta s_h$.

Proposition 1 (Feasible attenuation band). *Under Definition 4, assume the soft-injection case $s_a \leq \beta s_h$. Then (i) availability against the measured distractor requires $\beta > \rho$; (ii) rank preservation of the direct hit over the injected sibling requires $\beta < 1$, and more generally $s'_a < s_h$; (iii) when $\rho < 1$, any $\beta \in (\rho, 1)$ satisfies both conditions in the active injection case. If $s_a > s_h$ already, rank preservation depends on the observed sibling score rather than on β alone.*

Proof. If $s_a \leq \beta s_h$, then $s'_a = \beta s_h$. Availability follows from $\beta s_h > s_d$, i.e. $\beta > \rho$. Rank preservation requires $s'_a < s_h$, which reduces to $\beta < 1$ when $s_h > 0$. $\square$

Exploratory soft-rank diagnostic. We also minimize an exploratory soft pairwise diagnostic on a validation grid,

$$\mathcal{L}_{\text{sr}}(\beta; \rho) = -\log \sigma(\tau(\beta - \rho)) - \gamma \log \sigma(\tau(1 - \beta)), \quad (5)$$

with logistic σ , temperature τ , and rank weight γ . This visualization is useful for studying the availability–ranking trade-off, but it depends on the diagnostic pool, query channel, and chosen hyperparameters. Default $\beta=0.98$ is selected from the sweep in Table A.6, where $\{0.90, 0.95, 0.98\}$ form a stable high-AH band.

Fair controls include hard parent hydration, session-max aggregation, and shuffled *sid*. On two-turn LegalEp episodes, hard hydration and session-max coincide under full-sibling expansion; we report a single hard-expansion

Algorithm 2 Gated Hybrid Soft O2 + Soft O2-C

Require: Base hits $\mathcal{R}_0$, budgets (k, B_c) , attenuations (β, β_c)

```
1:  $\mathcal{R} \leftarrow \text{SoftBind}(\mathcal{R}_0; \text{Sib}, \beta)$  ▷ session pathway first
2:  $\mathcal{S}_{\text{cov}} \leftarrow \{sid_t : t \in \text{top-}k(\mathcal{R})\}$ 
3:  $\mathcal{T}_{\text{trig}} \leftarrow \{t \in \mathcal{R}_0\}$  ▷ direct dense hits only
4: for each trigger  $t \in \mathcal{T}_{\text{trig}}$  do
5:    $s_{\text{max}} \leftarrow \max\{s_u : u \in \mathcal{R}, cid_u = cid_t\}$ 
6:   for each  $t' \in \text{Clus}(t)$  with  $sid_{t'} \notin \mathcal{S}_{\text{cov}}$ , up to budget  $B_c$  do
7:      $s_{t'} \leftarrow \max(s_{t'}, \beta_c \cdot s_{\text{max}})$ 
8:   end for
9: end for
10: return top- $k$  of  $\mathcal{R}$ 
```

254 baseline. Indexes that concatenate an entire question–answer episode into
255 one retrieval unit (BM25-joint or dense joint-episode documents) change the
256 input granularity and are therefore outside the Soft O2 turn-level protocol;
257 we do not report them as Soft O2 competitors.

258 3.5. Semantic consolidation and Soft O2-C

259 Semantic consolidation follows a BIRCH-style dual-phase procedure Zhang
260 et al. (1996).

261 **Definition 5** (BIRCH dual-phase consolidation). *Let $\{\mathbf{c}_j\}$ be current cen-*
262 *troids and θ_H, θ_M the assign /merge thresholds. **Phase 1 (assign).** For*
263 *a new embedding $\mathbf{v}$, set $j^* = \arg \max_j S(\mathbf{v}, \mathbf{c}_j)$; if $S(\mathbf{v}, \mathbf{c}_{j^*}) \geq \theta_H$ then*
264 *$C_{j^*} \leftarrow C_{j^*} \cup \{\mathbf{v}\}$ and refresh $\mathbf{c}_{j^*}$, else spawn a new cluster. **Phase 2***
265 *(merge). Periodically, merge pairs with $S(\mathbf{c}_i, \mathbf{c}_j) \geq \theta_M$ and purge near-*
266 *duplicates above a redundancy threshold. Under O3, both phases run inside*
267 **Finalize.** *Each turn t receives a cluster id cid_t , inducing $\text{Clus}(t)$ in Defini-*
268 *tion 1.*

269 **Definition 6** (Soft O2-C). *Soft O2-C is $\text{SoftBind}(\cdot; \text{Clus}, \beta_c)$ with default*
270 *$\beta_c=0.95 < \beta$, reflecting noisier thematic neighbourhoods than session mem-*
271 *bership. A per-cluster sibling budget B_c caps $|\text{Clus}(t) \cap \text{Injected}|$ so large*
272 *thematic clusters cannot flood top- k .*

273 Gated Hybrid (Algorithm 2) injects only *cross-session* cluster siblings
274 relative to sessions already covered by dense /Soft O2 hits, and triggers




Figure 2: Shared-corpus protocol. Gold needles and same-domain distractors share one store; queries probe multiple channels with turn-level systems and significance tests. Exact replay is diagnostic; paraphrase, follow-up, and advice-recall are primary channels.

cluster binding only on direct dense hits. This prevents thematic confusers from displacing same-session siblings already recovered by Soft O2. Soft O2 and Soft O2-C are complementary: Soft O2 is the default binder when query and gold share *sid*; Soft O2-C is required when gold evidence is held in a different session but the same cluster.

4. Corpora and Evaluation Protocol

Figure 2 sketches the shared-corpus design used throughout the legal evaluation. Gold needles and same-domain distractors occupy one store. Each query probes that shared archive under fixed interference conditions. Privately rebuilt haystacks are avoided so that system comparisons isolate the retrieval operator.

4.1. Datasets and query channels

Multi-turn authenticity is evaluated on CAIL2024 consultation dialogues. Gold needles are authentic preliminary and final dialogues, with conversation turns completed by label turns where applicable. Query channels cover the first user turn ($U1$), a later in-dialogue user turn ($U_k\text{-followup}$), the last user turn ($U\text{-last}$), and constrained paraphrases ($U\text{-para}$). Evidence defaults to assistant turns in the target session.

LegalEp-DISC and LegalEp-Lawyer rebuild public DISC-Law-SFT and Lawyer-LLaMA subsets into consultation-episode corpora. Each episode is a natural (question, answer) pair with session identifier *sid*. Reconstruction applies length bounds and legal-content filters, excludes exam prompts, removes near-duplicates, separates needles from distractors under a fixed seed, and retains an audit subsample. After filtering, each LegalEp corpus retains 3500 passed sessions (500 needles, 2500 distractors at the $M=3000$ primary tier).

Query-channel policy (addressing exact-match shortcuts).. Exact replay uses the stored user question and is reported only as a *diagnostic* upper bound on surface-form recovery. Primary claims use paraphrase, Uk-followup /U-last, and advice-recall queries whose surface form is not identical to the stored question text; paraphrase and advice-recall templates are generated under a fixed seed and audited on a subsample. LegalMem-MT reuses CAIL multi-turn gold under the same store for same-session Soft O2 vs Soft O2-C contrasts and for Mix reconstructions.

Archive size is varied with gold needles held fixed. Unless stated otherwise, Soft O2 transfer tables use the $M \approx 3000$ tier; O1–O2 ablations use $M=400$, $S=300$.

4.2. Metrics and baselines

Primary metrics are Answer Hit@ k (AH), Episode Completeness@ k (EC), and corrected normalized Discounted Cumulative Gain@ k (nDCG) at $k=10$. AH is binary: whether a gold answer turn appears in the top- k . EC is the mean coverage of gold turns within the target session. Session Hit@ k is reported only for diagnosis. For nDCG we use fixed graded relevance over the full gold session: answer turns $rel=1.0$, other same-session turns $rel=0.5$, all other turns $rel=0.0$. The ideal denominator (IDCG) is computed from all gold-relevant turns in the target session and truncated at k , so every system on a query shares the same IDCG. Earlier draft values based on retrieved-item IDCG are not comparable. We therefore treat Answer Hit and Episode Completeness as the primary numeric claims in the main Soft O2 grids, report a failure taxonomy derived from Session Hit versus Answer Hit, and release a corrected-nDCG recompute script (`paper/scripts/recompute_corrected_metrics.py`) with fixed-gold IDCG artifacts under `paper/ipm/figures/`. Primary contrasts use bootstrap 95% CIs and McNemar mid- p tests on paired per-query

hits; where families of related hypotheses are interpreted jointly, we also report Holm-corrected p -values (Table A.7).

Turn-level dense FlatIP (O1) is the base dense retriever. Soft O2 applies soft sibling score inheritance on the episodic pathway. Soft O2-C applies the analogous rule within BIRCH clusters. Hard parent hydration expands a hit to all siblings by copying the trigger score. Lexical controls include BM25 over turns and dense+BM25 Reciprocal Rank Fusion (RRF) under the same turn-level evidence unit. We do *not* include BM25-joint or dense joint-episode indexes in Soft O2 primary comparisons: those methods retrieve pre-concatenated episode documents and therefore answer a different retrieval problem. A cross-encoder (CE) control reranks the top-30 FlatIP candidates with BAAI/bge-reranker-v2-m3. Negative and sensitivity controls include shuffled *sid*, a β sweep, and IVFPQ as a quantization reference.

4.3. Fair Mix protocol for Soft O2-C

On fully same-session LegalMem grids, Soft O2 already recovers gold answers that share *sid* with the query, so Soft O2-C has little exclusive headroom and may inject thematic confusers. We therefore reconstruct the *same* LegalEp and LegalMem archives under a Mix protocol: 70% of gold episodes are split into query-side S_a and evidence-side S_b with distinct session identifiers, while 30% remain intact. Distractors are unchanged. All operators—FlatIP, Soft O2, Soft O2-C, and gated Hybrid—share unrestricted dense retrieval on this store; we report overall AH and stratum AH on `cross_session` vs `same_session` queries.

Glue as a solvability bound (not a retrieval cheat).. After BIRCH consolidation on the Mix store, each cross-session S_a/S_b split pair is *glued* into one cluster identifier. This step uses gold split metadata available only in the evaluation protocol—it does *not* inject labels at query time and is *not* used in real unsupervised clustering deployments. Its sole purpose is to establish a **solvability bound**: if Soft O2-C cannot recover cross-session gold when the semantic store *must* contain both halves of a split pair in one cluster, then failure reflects mechanism limits rather than accidental cluster separation. Glue therefore tests whether cluster soft binding *can* complete episodes when the slow pathway has the requisite thematic co-membership; natural same-cluster rates without glue are lower and are reported only as diagnostics (Section 6.3). Soft O2’s session membership constraint is unchanged—glue affects only whether Soft O2-C has a cluster in which to bind, not

whether Soft O2 receives privileged label access. LegalEp Mix uses advice-recall queries; LegalMem Mix uses mid-split Uk-followup queries. External LongMemEval /LoCoMo numbers from prior retrieval runs on this codebase (QA correctness 70.7% / 68.2% on released test splits) are cited only as long-horizon memory context Wu et al. (2024); Maharana et al. (2024), not as Soft O2-C evidence on legal corpora.

5. Experiments

Unless noted, the embedding model is `nomic-embed-text-v2-moe` Nussbaum et al. (2025), $k=10$, seed 42. We report primary contrasts in figures and keep full numeric grids in the appendix. Result files and scripts are released with the accompanying repository so that each table maps to a public JSON artifact.

5.1. Implementation operators and Soft O2 ablation

Table 3 reports the shared-store ablation on DISC-Law and Lawyer-LLaMA ($M=400$, $S=300$). IVFPQ baseline Answer Hit is 0.540 / 0.397. O1 alone yields a small gain on DISC (0.567) and a near-tie on Lawyer (0.397), confirming that exact indexing is a necessary implementation prerequisite but insufficient alone. O1+Soft O2 raises Answer Hit to 0.780 / 0.680 (+0.240 / +0.283 over baseline) and Episode Completeness to 0.888 / 0.840—the primary algorithmic gain. O3 is enabled for all configurations in this ladder as a construction prerequisite; it does not alter retrieval scores.

Table 3: O1+Soft O2 ablation on the shared store ($M=400$, $S=300$, $k=10$). EC denotes Episode Completeness (gold-turn coverage); AH denotes Answer Hit. O3 bulk-load is on for all rows. Boldface: best AH and EC within each corpus. Corrected nDCG is reported for the $M \approx 3000$ grids (Table 5), not for this small-store ladder.

Corpus	System	EC@10	AH@10
DISC-Law	IVFPQ baseline	0.770	0.540
	+ O1 FlatIP	0.782	0.567
	+ O1+O2 Soft O2	0.888	0.780
Lawyer-LLaMA	IVFPQ baseline	0.698	0.397
	+ O1 FlatIP	0.698	0.397
	+ O1+O2 Soft O2	0.840	0.680

Table 4: RQ1 failure taxonomy ($k=10$) from the FlatIP rebuild with fixed-gold metrics. Incomplete = target session partially retrieved but answer missing. Soft O2 sharply reduces incompleteness while session-miss rates stay largely unchanged, motivating session soft binding before Soft O2-C.

Corpus / channel	System	Complete	Incomplete	Session miss
LegalEp-DISC / u-para	FlatIP	70.8%	26.8%	1.6%
LegalEp-DISC / u-para	Soft O2	89.1%	8.7%	1.8%
LegalEp-DISC / advice	FlatIP	48.0%	17.0%	35.0%
LegalEp-DISC / advice	Soft O2	60.8%	3.8%	35.4%
LegalEp-Lawyer / u-para	FlatIP	71.5%	25.3%	0.8%
LegalEp-Lawyer / u-para	Soft O2	94.0%	3.8%	1.6%
CAIL / Uk-followup	FlatIP	45.8%	54.0%	0.2%
CAIL / Uk-followup	Soft O2	92.8%	7.0%	0.2%

385 5.2. RQ1 failure taxonomy

386 We classify each query into *complete* (answer evidence in top- k), *incom-*
387 *plete* (target session partially retrieved but answer missing), or *session miss*
388 (no target-session turn in top- k). Table 4 shows that incomplete retrieval is a
389 substantial recurrent failure mode for FlatIP on primary non-exact channels,
390 motivating Soft O2.

391 5.3. Turn-level lexical controls

392 Primary Soft O2 claims stay within a turn-level index: FlatIP, Soft O2,
393 hard hydration, and shuffled *sid* share the same episodic store and evidence
394 unit. As a same-granularity lexical control, BM25-turn indexes each turn
395 independently (Table A.4). On CAIL non-exact channels Soft O2 exceeds
396 BM25-turn (Uk 0.698 vs 0.482; U-last 0.762 vs 0.413), while BM25-turn
397 remains competitive on exact replay where surface overlap is high. We de-
398 liberately omit BM25-joint and dense joint-episode indexing from Soft O2
399 head-to-head tables: both concatenate question and answer into one docu-
400 ment before retrieval, so a hit already contains the answer and the comparison
401 would confound Soft O2’s binding mechanism with a change of input gran-
402 ularity. Those coarser-grained designs are left as future work rather than
403 Soft O2 baselines.

404 5.4. Soft O2 on session (CAIL and LegalEp)

405 Figure 3 summarizes published V4 AH@10 and EC@10 on authentic CAIL
406 multi-turn channels at $M \approx 3000$ under O1 FlatIP—these absolute AH lev-

Table 5: Corrected graded nDCG@10 (fixed-gold IDCG) under a fresh FlatIP rebuild, distinct from the published V4 AH/EC grids in Tables A.1–A.3. Boldface: Soft O2 better than FlatIP within each corpus/channel pair on AH and nDCG. Soft O2 reduces incompleteness while raising both metrics.

Corpus	Channel	System	AH@10	nDCG@10	Incomplete
CAIL	U1	FlatIP	0.832	0.554	16.8%
CAIL	U1	Soft O2	0.960	0.776	4.0%
CAIL	Uk	FlatIP	0.458	0.445	54.0%
CAIL	Uk	Soft O2	0.928	0.793	7.0%
CAIL	U-last	FlatIP	0.425	0.420	57.3%
CAIL	U-last	Soft O2	0.945	0.791	5.3%
LegalEp-DISC	exact	FlatIP	0.742	0.730	25.8%
LegalEp-DISC	exact	Soft O2	0.924	0.813	7.6%
LegalEp-DISC	u-para	FlatIP	0.716	0.681	26.8%
LegalEp-DISC	u-para	Soft O2	0.895	0.759	8.7%
LegalEp-DISC	advice	FlatIP	0.480	0.466	17.0%
LegalEp-DISC	advice	Soft O2	0.608	0.524	3.8%
LegalEp-Lawyer	exact	FlatIP	0.732	0.753	26.8%
LegalEp-Lawyer	exact	Soft O2	0.976	0.869	2.4%
LegalEp-Lawyer	u-para	FlatIP	0.739	0.744	25.3%
LegalEp-Lawyer	u-para	Soft O2	0.946	0.844	3.8%
LegalEp-Lawyer	advice	FlatIP	0.444	0.483	23.4%
LegalEp-Lawyer	advice	Soft O2	0.652	0.582	2.2%

407 els are the Soft O2 transfer claims used throughout the paper. Soft O2 lifts
 408 Answer Hit over FlatIP on every channel (U1 0.788 vs 0.570**; Uk 0.698 vs
 409 0.270**; U-last 0.762 vs 0.245**). Shuffled *sid* removes the gain (U1 0.452; Uk
 410 0.230; U-last 0.200), tying the improvement to correct episode membership
 411 (RQ1–RQ2). In the same V4 AH/EC grids, hard expansion can raise AH
 412 slightly above Soft O2 (e.g., U1 0.802 vs 0.788) while lowering EC (U1 0.578
 413 vs Soft O2 0.664). Soft O2 therefore behaves as a rank-preserving binder: sib-
 414 lings enter under attenuated scores, and episode coverage remains higher than
 415 under hard score copy. Under corrected fixed-gold IDCG on a fresh FlatIP
 416 rebuild (Table 5; absolute AH levels differ from V4), Soft O2 also raises
 417 nDCG over FlatIP on every CAIL channel (U1 0.776 vs 0.554; Uk 0.793 vs
 418 0.445; U-last 0.791 vs 0.420) while cutting incompleteness sharply (Table 4).
 419 Soft O2 still exceeds FlatIP+CE on every CAIL channel (Table A.5). Full
 420 V4 rows with corrected nDCG appear in Table A.1.

421 Figure 4 reports LegalEp transfer on the same published V4 AH grids.

CAIL2024 multi-turn results ($M \approx 3000$, turn-level systems)




Figure 3: CAIL2024 multi-turn results at tier $M \approx 3000$ (published V4 AH/EC). Soft O2 improves AH and EC over FlatIP; shuffled *sid* removes the gain; hard expansion can maximize AH at a modest EC cost.

On LegalEp-DISC exact replay Soft O2 and FlatIP stay near-tied (0.638 vs 0.640), where both turns are already recoverable for many queries—consistent with treating exact as diagnostic. On Lawyer exact Soft O2 improves AH to 0.752 from FlatIP 0.692 (McNemar mid- $p=0.024^*$). Paraphrase yields the clearest LegalEp gains: Soft O2 reaches 0.648 vs FlatIP 0.547 on DISC and 0.759 vs 0.568 on Lawyer, exceeding hard expansion on both corpora (mid- $p<0.001^{**}$ on both corpora). Advice-recall queries omit the stored question surface form: Soft O2 raises AH to 0.428 from FlatIP 0.342 on DISC and to 0.472 from 0.332 on Lawyer, again above hard expansion (mid- $p<0.001^{**}$), while shuffled *sid* falls toward or below FlatIP. Corrected rebuild nDCG mirrors the same Soft O2>FlatIP ordering on paraphrase and advice-recall (Table 5). Detailed grids are in Tables A.2–A.3.

5.5. Same-store cluster pathway validation

Table 6 contrasts Soft O2 and ungated Soft O2-C on standard same-session LegalEp /LegalMem stores (unrestricted dense; gold answers share *sid* with the query). Soft O2 dominates Soft O2-C on every reported channel (e.g., LegalMem U1 AH 0.903 vs 0.712; LegalEp-DISC exact 0.836 vs 0.744). The slow pathway is therefore *not* a free upgrade when episode membership already encodes the gold answer: cluster siblings inject thematic confusers that displace session-complete evidence.

Table 7 reports the complementary fair Mix setting (RQ3–RQ4): the same stores with 70% cross-session Sa/Sb gold and 30% intact same-session gold, unrestricted dense for every operator. Gated Hybrid Soft O2+Soft O2-C

LegalEp ($M=3000$): Soft O2 gains are strongest on paraphrase and remain positive on advice-recall




Figure 4: LegalEp AH@10 ($M=3000$, 500 needles). Soft O2 gains are strongest on paraphrase and remain positive on advice-recall; exact transfer is corpus-dependent and treated as diagnostic.

Table 6: Same-session Soft O2 vs Soft O2-C on the shared LegalEp /LegalMem stores (AH@10). Soft O2-C is ungated cluster soft binding. Boldface: better AH within each row.

Corpus / channel	Soft O2	Soft O2-C
LegalEp-DISC / exact	0.836	0.744
LegalEp-DISC / u-para	0.807	0.730
LegalEp-DISC / advice	0.548	0.490
LegalEp-Lawyer / exact	0.818	0.710
LegalEp-Lawyer / u-para	0.817	0.733
LegalEp-Lawyer / advice	0.546	0.448
LegalMem-MT / U1	0.903	0.712
LegalMem-MT / Uk	0.793	0.385

445 improves over Soft O2 on LegalEp-Lawyer overall (AH 0.534 vs 0.496; mid-
 446 $p=0.010^*$) and on the cross-session stratum (0.523 vs 0.446; mid- $p<10^{-4**}$).
 447 On LegalEp-DISC, Hybrid is directionally above Soft O2 (0.486 vs 0.482)
 448 but not significant at $N=500$. On LegalMem-MT, Hybrid attains the highest
 449 overall AH (0.646 vs Soft O2 0.618; mid- $p=0.18$; cross-session mid- $p=0.095$)
 450 while preserving most same-session hits (0.867 vs 0.887). Ungated Soft O2-
 451 C can raise cross-session hits but often displaces same-session siblings; Hy-
 452 brid triggers cluster binding only on direct dense hits. Together, Tables 6–7
 453 support a staged narrative: Soft O2 is the default same-session operator;
 454 Soft O2-C /Hybrid are exploratory complements that can recover additional
 455 cross-session gold under the Mix solvability-bound protocol while Soft O2

remains strong on intact same-session gold.

Table 7: Fair Mix protocol (same store, unrestricted dense; AH@10). Gold is a 70%/30% mix of cross-session Sa/Sb splits and intact same-session episodes. Soft O2-C denotes ungated cluster binding; Hybrid is Soft O2+gated Soft O2-C. Boldface: best overall AH within each corpus (and best cross-session AH when Hybrid uniquely leads that stratum). Significance vs Soft O2: * $p<0.05$, ** $p<0.01$.

Corpus	System	AH	AH _{cross}	AH _{same}	mid- p vs Soft O2
LegalEp-DISC Mix	FlatIP	0.478	0.477	0.480	—
	Soft O2	0.482	0.443	0.573	—
	Soft O2-C	0.410	0.443	0.333	0.0004**
	Hybrid	0.486	0.469	0.527	0.8099
LegalEp-Lawyer Mix	FlatIP	0.462	0.454	0.480	—
	Soft O2	0.496	0.446	0.613	—
	Soft O2-C	0.464	0.491	0.400	0.0857
	Hybrid	0.534*	0.523	0.560	0.0105*
LegalMem-MT Mix	FlatIP	0.534	0.526	0.553	—
	Soft O2	0.618	0.503	0.887	—
	Soft O2-C	0.500	0.537	0.413	0.0000**
	Hybrid	0.646	0.551	0.867	0.1837

McNemar mid- p tests on paired Answer Hit; * $p<0.05$, ** $p<0.01$.

5.6. Controls, significance, scale, and QA audit

Lexical BM25-turn is competitive on exact replay (DISC exact 0.802) and weaker on advice-recall (0.382). Dense+BM25 RRF complements Soft O2 on surface-overlap channels under the same turn-level unit. Cross-encoder reranking (`bge-reranker-v2-m3`, top-30) improves FlatIP inside the CE suite, yet remains below Soft O2 on CAIL (Uk FlatIP+CE 0.375 vs Soft O2 0.698; Table A.5). A β sweep peaks near $\{0.9, 0.95, 0.98\}$ and weakens at $\beta=0.5$ (Table A.6), consistent with the diagnostic gap-ratio band in Section 3.4 (Table 2). McNemar mid- p tests (Table A.7) show Soft O2 versus FlatIP highly significant on all CAIL channels and on LegalEp paraphrase /advice-recall.

Figure 5 and Table A.8 report quality and warm-query latency on LegalEp-DISC exact ($Q=100$, three repeats; $M \in \{100, 400, 1600, 6400\}$). Soft O2 retains an AH advantage at each tier; IVFPQ collapses at $M=1600$ (AH 0.330), motivating O1; absolute latency grows with M . We treat O1 FlatIP




Figure 5: Scale curve on DISC-Law exact. Soft O2 keeps an AH edge over FlatIP as M grows; warm-query latency scales with archive size; IVFPQ quality collapses at moderate M .

as an engineering prerequisite for Soft O2 at these scales and avoid claiming deployment readiness beyond the measured curve.

End-to-end QA audit (summary).. A dual-protocol QA audit ($N=270$ per corpus; generator `qwen3:14b`, independent judge `qwen3:32b`) links retrieval quality to downstream answerability under a *single* retrieval configuration (Soft O2); we do not claim causal superiority over FlatIP in generation without paired downstream controls. Full protocol and stratified results appear in Section 6.2. On an older revision-protocol store ($M=400$, five seeds), Soft O2 AH on LegalEp-DISC is 0.851 ± 0.008 versus FlatIP 0.698 ± 0.040 .

6. Discussion

6.1. Findings and scope

RQ1–RQ2 show that episode incompleteness is a substantial recurrent FlatIP failure mode under same-domain interference, and that Soft O2 converts many incomplete retrievals into complete ones without shrinking session-miss rates (Table 4). O1 restores score fidelity as an implementation prerequisite; Soft O2 supplies the primary algorithmic gain through soft sibling inheritance; O3 enables archive construction at scale. The diagnostic gap-ratio analysis (Section 3.4) explains why β should remain close to but below unity, and why the validation sweep peaks in $\{0.9, 0.95, 0.98\}$ rather than at hard hydration ($\beta=1$). On CAIL, Soft O2 raises Answer Hit by large margins in the published V4 grids (U1 +0.218, Uk +0.428, U-last +0.517)

with collapsing shuffled-*sid* controls, and keeps higher EC than hard expansion despite occasional AH ties. LegalEp paraphrase and advice-recall—not exact replay—carry the primary Soft O2 transfer claims. Corrected fixed-gold nDCG on a FlatIP rebuild corroborates the Soft O2>FlatIP ordering without replacing the V4 AH/EC transfer numbers (Table 5). RQ3 shows that the BIRCH slow pathway benefits Soft O2-C only when gold is not already co-session; Mix Hybrid Soft O2+Soft O2-C recovers that regime without forcing Soft O2 to zero on intact same-session gold. RQ4: Soft O2 (session) is the default operator; Soft O2-C /Hybrid (cluster) are complementary exploratory bounds within one BIMS-LEGAL dual-store. Sparse, fusion, and cross-encoder baselines situate Soft O2 among standard IR alternatives Robertson and Zaragoza (2009); Cormack et al. (2009); Nogueira et al. (2019).

6.2. RAG application perspective

Consultation-memory retrieval is judged downstream by whether an LLM can answer from recovered evidence rather than confabulate. We audit this end-to-end with a dual-protocol pipeline on Soft O2 retrieval only: generator **qwen3:14b** and independent judge **qwen3:32b** are distinct models. Per corpus we pool $N=270$ judged instances as P1 $n=120$ plus P2 $n=150$ (folder **legal_qa_n270**). Table 8 summarizes pooled Evidence-Answerability Rate (EAR) and stratified independent-judge scores. When Soft O2 retrieves complete episodes—both question and answer turns in top- k —the generator receives grounded prior advice and judged answerability rises in this single-system audit (pooled EAR 0.867 on DISC and 0.859 on Lawyer; Wilson 95% CI [0.83, 0.90] / [0.82, 0.89]). Follow-up queries remain the hardest stratum (independent-judge EAR 0.517 / 0.480), consistent with episodic incompleteness persisting even when surface-form overlap is low.

A blind two-annotator legal audit ($n=60$ /corpus) yields Cohen’s $\kappa=0.71$ / 0.68 and separates retrieval from generation failures (retrieval-failure share 0.283 / 0.317). Wrong-session contamination rates stay low (0.083 / 0.100), suggesting that complete episode retrieval under Soft O2 is associated with lower rates of a common RAG failure mode in which the model answers from a neighbouring client’s matter. Hallucinated-legal-basis rates (0.117 / 0.133) remain non-zero, indicating that grounding in retrieved episodes does not guarantee statutory correctness—a scope boundary for deployment. AH and EAR measure evidence availability and judged answerability respectively; statutory correctness, citation validity, and professional

liability lie outside this evaluation.

Table 8: End-to-end QA audit ($N=270$ per corpus; generator `qwen3:14b`, independent judge `qwen3:32b`). EAR: Evidence-Answerability Rate (judge score ≥ 0.7). Stratified judge scores: $n=90$ /corpus across exact, paraphrase, and follow-up protocols.

Corpus	Pooled EAR	Stratified judge	Human κ
DISC-Law	0.867	0.757	0.71
Lawyer-LLaMA	0.859	0.729	0.68

6.3. Limitations

LegalEp episodes are consultation pairs; multi-turn authenticity remains reserved for CAIL. Mix glue establishes a solvability bound for Soft O2-C evaluation (Section 4.3); natural same-cluster rates without glue are lower and are diagnostics only—real unsupervised BIRCH clustering would not receive split-pair metadata at build time. Soft O2 depends on correct session identifiers; Soft O2-C depends on cluster quality. The β diagnostic is semi-empirical: it assumes a dominant distractor and a well-defined hit score, and the fitted ρ uses a stored user-turn embedding as an exact-channel query proxy; paraphrase /advice-recall gap laws may shift the feasible band slightly but leave the high- β operating region unchanged. LLM paraphrases need human audit. Statutory correctness, citation validity, and professional liability remain unevaluated. Multiple channels inflate comparison multiplicity; we therefore report Holm-corrected p -values for the primary Soft O2 vs FlatIP family (Table A.7) alongside uncorrected McNemar mid- p tests.

Latency and deployment scope. Scale curves cover $M \in \{100, 400, 1600, 6400\}$ under a fixed hardware stack (Table A.8). At $M=6400$, warm-query p50 latency exceeds 4s (4061ms for Soft O2; 4458ms for FlatIP)—unacceptable for millisecond interactive search. BIMS-LEGAL should therefore be positioned as **high-precision offline or quasi-real-time consultation memory recovery**—e.g., pre-meeting case review, audit replay, or batch RAG over a firm’s advice archive—not as a sub-millisecond search engine. Future work will combine approximate nearest-neighbour indexes (e.g., HNSW) with Soft O2 and Soft O2-C binding layered on top of faithful candidate retrieval, preserving episode completion while reducing query latency. Behavior at roughly 10^5 sessions is

not measured, and we do not claim asymptotic scalability from O3 wall-clock alone.

7. Conclusion

BIMS-LEGAL targets consultation-memory retrieval under same-domain interference in Chinese legal archives through a dual-store design: an episodic turn index paired with a BIRCH semantic store. Soft O2 session soft binding is the primary algorithmic contribution; Soft O2-C with gated Hybrid is an exploratory cluster-binding complement under a gold-assisted Mix solvability bound, with Soft O2 attenuation β validated on a diagnostic gap-ratio model; O1 exact indexing and O3 bulk-load consolidation are implementation operators that enable the shared-store ablations reported here. Under a shared-corpus protocol spanning CAIL multi-turn dialogues and rebuilt LegalEp episode corpora, Soft O2 improves Answer Hit over FlatIP on multi-turn, paraphrase, and advice-recall channels with collapsing shuffled-*sid* controls, yields higher episode completeness than hard expansion on CAIL, and raises corrected fixed-gold nDCG while reducing incompleteness. Same-store Soft O2-C underperforms Soft O2 when gold shares *sid*; under fair Mix reconstructions, gated Hybrid Soft O2+Soft O2-C can recover additional cross-session gold while Soft O2 remains strong on intact same-session gold. End-to-end QA audit links complete episode retrieval to judged answerability and low wrong-session contamination in downstream generation under Soft O2 retrieval. Sparse, fusion, and cross-encoder baselines situate Soft O2 among standard turn-level IR alternatives; joint episode / document-level indexes are a different input granularity and are not Soft O2 competitors in this protocol.

Declarations

Funding

Omitted for double-blind review.

Conflict of interest

The authors declare no competing interests.

588 *Ethics approval*

589 This study uses publicly released research corpora and excludes confiden-
590 tial client data. The system is a research prototype for consultation-memory
591 retrieval.

592 *Data availability*

593 Code, processed evaluation corpora, and primary result files are pub-
594 licly available at [anonymizedrepositoryURL]. Upstream legal corpora:
595 CAIL2024 consultation tracks; Hugging Face ShengbinYue/DISC-Law-SFT
596 and Skepsun/lawyer_llama_data.

597 *Code availability*

598 The BIMS-LEGAL implementation (O1 FlatIP, Soft O2, O3 bulk-load,
599 Soft O2-C), LegalMem/LegalEp pipelines, Mix builders, and table/figure
600 scripts are released in the same repository with JSON artifacts for each main
601 table.

602 *Author contributions*

603 Omitted for double-blind review.

604 **References**

- 605 Ashley, K.D., 2017. Artificial Intelligence and Legal Analytics: New Tools
606 for Law Practice in the Digital Age. Cambridge University Press.
- 607 Askari, A., Aliannejadi, M., Abolghasemi, A., Kanoulas, E., Verberne, S.,
608 2024. Answer retrieval in legal community question answering, in: Ad-
609 vances in Information Retrieval (ECIR), Springer. pp. 477–485.
- 610 Atkinson, R.C., Shiffrin, R.M., 1971. The control of short-term memory.
611 Scientific American 225, 82–91.
- 612 Callan, J.P., Lu, Z., Xu, S., 1994. Passage-level evidence in document re-
613 trieval, in: Proceedings of the 17th Annual International ACM SIGIR
614 Conference on Research and Development in Information Retrieval, pp.
615 302–310.
- 616 Chen, J., Xiao, S., Zhang, P., Luo, K., Lian, D., Liu, Z., 2024. M3-
617 embedding: Multi-linguality, multi-functionality, multi-granularity text
618 embeddings through self-knowledge distillation, in: Findings of the associ-
619 ation for computational linguistics: ACL 2024, pp. 2318–2335.
- 620 Cormack, G.V., Clarke, C.L.A., Buettcher, S., 2009. Reciprocal rank fusion
621 outperforms condorcet and individual rank learning methods, in: Proceed-
622 ings of the 32nd International ACM SIGIR Conference on Research and
623 Development in Information Retrieval, pp. 758–759.
- 624 Dai, Z., Xiong, C., Liu, Z., Callan, J., 2019. Context-aware sentence/passage
625 encoding for answering factoid questions, in: Proceedings of the 2019 Con-
626 ference on Empirical Methods in Natural Language Processing and the
627 9th International Joint Conference on Natural Language Processing, pp.
628 241–251.
- 629 Gutiérrez, B.J., Shu, Y., Gu, Y., Yasunaga, M., Su, Y., 2024. Hipporag:
630 Neurobiologically inspired long-term memory for large language models.
631 Advances in neural information processing systems 37, 59532–59569.
- 632 Huang, Q., Tao, M., Zhang, C., An, Z., Jiang, C., Chen, Z., Wu, Z., Feng,
633 Y., 2023. Lawyer llama technical report. [arXiv:2305.15062](https://arxiv.org/abs/2305.15062).
- 634 Johnson, J., Douze, M., Jégou, H., 2019. Billion-scale similarity search with
635 GPUs. IEEE Transactions on Big Data 7, 535–547.

- 636 Karpukhin, V., Oguz, B., Min, S., Lewis, P., Wu, L., Edunov, S., Chen,
637 D., Yih, W.t., 2020. Dense passage retrieval for open-domain question
638 answering, in: Proceedings of the 2020 Conference on Empirical Methods
639 in Natural Language Processing (EMNLP), pp. 6769–6781.
- 640 Kumaran, D., Hassabis, D., McClelland, J.L., 2016. What learning systems
641 do intelligent agents need? complementary learning systems theory up-
642 dated. Trends in cognitive sciences 20, 512–534.
- 643 Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küt-
644 tler, H., Lewis, M., Yih, W.t., Rocktäschel, T., et al., 2020. Retrieval-
645 augmented generation for knowledge-intensive nlp tasks. Advances in neural
646 information processing systems 33, 9459–9474.
- 647 Liu, B., Wu, Y., Zhang, F., Liu, Y., Wang, Z., Li, C., Zhang, M., Ma, S., 2022.
648 Query generation and buffer mechanism: Towards a better conversational
649 agent for legal case retrieval. Information Processing & Management 59,
650 103051.
- 651 Louis, A., van Dijck, G., Spanakis, G., 2024. Interpretable long-form legal
652 question answering with retrieval-augmented large language models, in:
653 Proceedings of the AAAI Conference on Artificial Intelligence, pp. 22266–
654 22275.
- 655 Maharana, A., Lee, D.H., Tulyakov, S., Bansal, M., Barbieri, F., Fang, Y.,
656 2024. Evaluating very long-term conversational memory of llm agents, in:
657 Proceedings of the 62nd Annual Meeting of the Association for Computa-
658 tional Linguistics (Volume 1: Long Papers), pp. 13851–13870.
- 659 Martinez-Gil, J., 2023. A survey on legal question–answering systems. Com-
660 puter Science Review 48, 100552.
- 661 McClelland, J.L., McNaughton, B.L., O’Reilly, R.C., 1995. Why there are
662 complementary learning systems in the hippocampus and neocortex: in-
663 sights from the successes and failures of connectionist models of learning
664 and memory. Psychological review 102, 419.
- 665 Mo, F., Gao, Y., Wu, Z., Liu, X., Chen, P., Li, Z., Wang, Z., Li, X., Jiang,
666 M., Nie, J.Y., 2026. Leveraging historical information to boost retrieval-
667 augmented generation in conversations. Information Processing & Man-
668 agement 63, 104449.

- 669 Modarressi, A., Imani, A., Fayyaz, M., Schütze, H., 2023. Ret-llm: Towards
670 a general read-write memory for large language models. arXiv preprint
671 arXiv:2305.14322 .
- 672 Nogueira, R., Yang, W., Cho, K., Lin, J., 2019. Multi-stage document rank-
673 ing with BERT. arXiv:1910.14497.
- 674 Nussbaum, Z., et al., 2025. Nomic embed multimodal: Interleaved text, im-
675 age, and datasets for general-purpose vision-language embeddings. Model
676 card: **nomic-embed-text-v2-moe** (Nomic AI). Embedding model used for
677 turn encoding in this work.
- 678 Packer, C., Wooders, S., Lin, K., Fang, V., Patil, S.G., Stoica, I., Gonzalez,
679 J.E., 2023. Memgpt: Towards llms as operating systems. arXiv preprint
680 arXiv:2310.08560 .
- 681 Rasmussen, P., Paliychuk, P., Beauvais, T., Ryan, J., Chalef, D., 2025. Zep:
682 A temporal knowledge graph architecture for agent memory. arXiv preprint
683 arXiv:2501.13956 .
- 684 Robertson, S., Zaragoza, H., 2009. The probabilistic relevance framework:
685 BM25 and beyond. Foundations and Trends in Information Retrieval 3,
686 333–389.
- 687 Tan, H., Zhang, Z., Ma, C., Chen, X., Dai, Q., Dong, Z., 2025. Mem-
688 bench: Towards more comprehensive evaluation on the memory of llm-
689 based agents, in: Findings of the Association for Computational Linguis-
690 tics: ACL 2025, pp. 19336–19352.
- 691 Tulving, E., et al., 1972. Episodic and semantic memory. Organization of
692 memory 1, 1.
- 693 Wang, Y., Gao, Y., Chen, X., Jiang, H., Li, S., Yang, J., Yin, Q., Li, Z.,
694 Li, X., Yin, B., et al., 2024. Memoryllm: Towards self-updatable large
695 language models. arXiv preprint arXiv:2402.04624 .
- 696 Wu, D., Wang, H., Yu, W., Zhang, Y., Chang, K.W., Yu, D., 2024. Long-
697 memeval: Benchmarking chat assistants on long-term interactive memory.
698 arXiv preprint arXiv:2410.10813 .

699 Wu, Y., Liang, S., Zhang, C., Wang, Y., Zhang, Y., Guo, H., Tang, R.,
700 Liu, Y., 2025. From human memory to ai memory: A survey on memory
701 mechanisms in the era of llms. [arXiv:2504.15965](#).

702 Yue, S., Chen, W., Wang, S., Li, B., Shen, C., Liu, S., Zhou, Y., Xiao, Y.,
703 Yun, S., Huang, X., Wei, Z., 2023. Disc-lawllm: Fine-tuning large language
704 models for intelligent legal services. [arXiv:2309.11325](#).

705 Yue, S., Liu, S., Zhou, Y., Shen, C., Wang, S., Xiao, Y., Li, B., Yun, S.,
706 Shen, X., Chen, W., et al., 2024. Lawllm: Intelligent legal system with
707 legal reasoning and verifiable retrieval, in: International Conference on
708 Database Systems for Advanced Applications, Springer. pp. 304–321.

709 Zaheer, M., Kottur, S., Ravanbakhsh, S., Poczos, B., Salakhutdinov, R.,
710 Smola, A.J., 2017. Deep sets, in: Advances in Neural Information Process-
711 ing Systems.

712 Zhang, T., Ramakrishnan, R., Livny, M., 1996. BIRCH: an efficient data
713 clustering method for very large databases, in: Proceedings of the 1996
714 ACM SIGMOD International Conference on Management of Data, pp. 103–
715 114.

716 Zhang, Y., Zhai, S., Meng, Y., Bi, S., Chen, Y., Qi, G., 2024. Event is
717 more valuable than you think: Improving the similar legal case retrieval
718 via event knowledge. *Information Processing & Management* 61, 103729.

719 Zhong, W., Guo, L., Gao, Q., Ye, H., Wang, Y., 2024. Memorybank: En-
720 hancing large language models with long-term memory, in: Proceedings of
721 the AAAI Conference on Artificial Intelligence, pp. 19724–19731.

722 Appendix A. Numeric grids and controls

723 Main Soft O2 claims are visualized in Figures 3–5; O1–O2 and Mix
 724 Soft O2-C tables appear in the main text. Tables below retain full Soft O2 nu-
 725 meric rows for reproducibility: published V4 AH/EC with corrected rebuild
 726 nDCG. Boldface follows AH/EC winners so that Soft O2’s completeness ad-
 727 vantage remains visible even when hard hydration leads AH.

Table A.1: CAIL2024 shared-corpus results ($M \approx 3000$). Boldface marks the best AH@10 and best EC@10 within each channel (ties bolded jointly). **: Soft O2 vs FlatIP, McNemar mid- $p < 0.01$. Published V4 AH/EC are the Soft O2 transfer claims; nDCG@10 is corrected fixed-gold IDCG from a FlatIP rebuild and is not used for winner bolding.

Channel	System	AH@10	EC@10	nDCG@10	95% CI (AH)
U1	FlatIP	0.570	0.413	0.554	[0.530,0.610]
	Soft O2	0.788	0.664	0.776	[0.757,0.820]
	Hard hydr.	0.802	0.578	0.864	[0.768,0.835]
	Shuffled O2	0.452	0.363	0.411	[0.410,0.492]
Uk-followup	FlatIP	0.270	0.347	0.445	[0.237,0.307]
	Soft O2	0.698	0.654	0.793	[0.660,0.735]
	Hard hydr.	0.720	0.585	0.847	[0.685,0.753]
	Shuffled O2	0.230	0.320	0.369	[0.198,0.262]
U-last	FlatIP	0.245	0.335	0.420	[0.212,0.280]
	Soft O2	0.762	0.659	0.791	[0.725,0.797]
	Hard hydr.	0.790	0.586	0.854	[0.757,0.822]
	Shuffled O2	0.200	0.315	0.348	[0.172,0.232]

728 [†]nDCG@10 from fixed-gold IDCG FlatIP rebuild
 729 (recompute_corrected_metrics.py); published V4 AH/EC unchanged. Soft O2
 730 leads EC; hard hydration can lead AH.

Table A.2: LegalEp-DISC ($M=3000$, 500 needles). Boldface: best AH@10 and best EC@10 within each channel. * $p<0.05$, ** $p<0.01$ (Soft O2 vs FlatIP, McNemar mid- p). Hard hydration coincides with session-max aggregation; only hard hydration is shown. Published V4 AH/EC; nDCG@10 from corrected FlatIP rebuild (not used for winner bolding).

Channel	System	AH@10	EC@10	nDCG@10
exact	FlatIP	0.640	0.741	0.730
	Soft O2	0.638	0.780	0.813
	Hard hydr.	0.614	0.598	0.886
	Shuffled O2	0.448	0.645	0.666
advice-recall	FlatIP	0.342	0.438	0.466
	Soft O2	0.428	0.503	0.524
	Hard hydr.	0.402	0.405	0.572
	Shuffled O2	0.310	0.421	0.436
u-param	FlatIP	0.547	0.681	0.681
	Soft O2	0.648	0.735	0.759
	Hard hydr.	0.592	0.571	0.827
	Shuffled O2	0.469	0.610	0.627

Table A.3: LegalEp-Lawyer ($M=3000$, 500 needles). Boldface: best AH@10 and best EC@10 within each channel. * $p<0.05$, ** $p<0.01$ (Soft O2 vs FlatIP). Hard hydration coincides with session-max aggregation; only hard hydration is shown. Published V4 AH/EC; nDCG@10 from corrected FlatIP rebuild (not used for winner bolding).

Channel	System	AH@10	EC@10	nDCG@10
exact	FlatIP	0.692	0.806	0.753
	Soft O2	0.752	0.848	0.869
	Hard hydr.	0.612	0.615	0.897
	Shuffled O2	0.570	0.747	0.705
advice-recall	FlatIP	0.332	0.464	0.483
	Soft O2	0.472	0.540	0.582
	Hard hydr.	0.394	0.403	0.601
	Shuffled O2	0.348	0.467	0.454
u-param	FlatIP	0.568	0.718	0.744
	Soft O2	0.759	0.819	0.844
	Hard hydr.	0.596	0.594	0.864
	Shuffled O2	0.592	0.712	0.705

Table A.4: Turn-level BM25 and dense+BM25 RRF (AH@10). Boldface: higher AH within each row among reported turn-level cells. BM25-joint / joint-episode document indexes are omitted: they change evidence granularity and were not run as Soft O2 protocol competitors.

Corpus	Channel	AH@10
LegalEp-DISC	exact (BM25-turn)	0.802
LegalEp-DISC	advice-recall (BM25-turn)	0.382
LegalEp-DISC	exact (dense+BM25 RRF)	0.780
LegalEp-DISC	advice (dense+BM25 RRF)	0.394
CAIL	U1 / Uk / U-last (BM25-turn)	0.748 / 0.482 / 0.413
LegalEp-Lawyer	exact / advice (BM25-turn)	0.784 / 0.460

Table A.5: Cross-encoder control (AH@10). FlatIP and FlatIP+CE are paired in the CE suite (**bge-reranker-v2-m3** over top-30 dense candidates). Boldface: best AH within each row. Soft O2 reproduces the primary evaluation; FlatIP here may differ slightly from primary FlatIP due to independent CE-suite re-indexing.

Corpus / channel	FlatIP	FlatIP+CE	Soft O2
CAIL / U1	0.585	0.688	0.788
CAIL / Uk	0.265	0.375	0.698
CAIL / U-last	0.228	0.342	0.762
LegalEp-DISC / exact	0.572	0.722	0.638
LegalEp-Lawyer / exact	0.636	0.708	0.752

Table A.6: Soft O2 AH@10 versus β on the primary shared-corpus evaluation. Peak performance lies in $\{0.9, 0.95, 0.98\}$; aggressive attenuation ($\beta=0.5$) weakens binding.

Corpus	0.5	0.7	0.9	0.95	0.98	1.0
CAIL / Uk	0.323	0.537	0.697	0.700	0.698	0.662
LegalEp-Lawyer / exact	0.608	0.654	0.754	0.754	0.752	0.752
LegalEp-Lawyer / u-para	0.631	0.637	0.717	0.745	0.759	0.749
LegalEp-DISC / u-para	0.551	0.551	0.618	0.648	0.648	0.642

Table A.7: McNemar mid- p on paired Answer Hit (Soft O2 vs FlatIP / hard hydration / FlatIP+CE). Holm-adjusted p is reported for the nine primary Soft O2 vs FlatIP contrasts.

Corpus / channel	Contrast	ΔAH	mid- p	Holm p
CAIL / Uk	O2 vs FlatIP	+0.428	8.99e-48	8.09e-47
CAIL / Uk	O2 vs Hard hydr.	-0.022	0.243	—
CAIL / U1	O2 vs FlatIP	+0.218	5.19e-18	4.67e-17
CAIL / U1	O2 vs Hard hydr.	-0.013	0.435	—
CAIL / U-last	O2 vs FlatIP	+0.517	5.51e-65	4.96e-64
CAIL / U-last	O2 vs Hard hydr.	-0.028	0.0982	—
CAIL / Uk	O2 vs FlatIP+CE	+0.323	2.50e-26	—
CAIL / U1	O2 vs FlatIP+CE	+0.100	3.51e-05	—
CAIL / U-last	O2 vs FlatIP+CE	+0.420	3.92e-41	—
LegalEp-DISC / exact	O2 vs FlatIP	-0.002	0.934	0.934
LegalEp-DISC / exact	O2 vs Hard hydr.	+0.024	0.302	—
LegalEp-Lawyer / exact	O2 vs FlatIP	+0.060	0.0239	0.191
LegalEp-Lawyer / exact	O2 vs Hard hydr.	+0.140	3.50e-10	—
LegalEp-DISC / u-para	O2 vs FlatIP	+0.101	7.95e-05	6.36e-04
LegalEp-DISC / u-para	O2 vs Hard hydr.	+0.056	0.035	—
LegalEp-Lawyer / u-para	O2 vs FlatIP	+0.191	2.10e-15	1.68e-14
LegalEp-Lawyer / u-para	O2 vs Hard hydr.	+0.163	7.57e-09	—
LegalEp-DISC / advice	O2 vs FlatIP	+0.086	1.25e-04	8.75e-04
LegalEp-DISC / advice	O2 vs Hard hydr.	+0.026	0.215	—
LegalEp-Lawyer / advice	O2 vs FlatIP	+0.140	1.18e-10	9.44e-10
LegalEp-Lawyer / advice	O2 vs Hard hydr.	+0.078	5.72e-04	—

Table A.8: Scale curve on DISC-Law (mean over 3 repeats). Latency is warm-query p50/p95. Boldface: best AH@10 within each M block; lowest p50/p95 within each M block.

M	Mode	AH@10	Build (s)	p50 (ms)	p95 (ms)
100	FlatIP	0.963	3.0	67	80
100	IVFPQ	0.613	25.2	40	54
100	Soft O2	0.970	3.2	63	69
400	FlatIP	0.890	26.5	248	305
400	IVFPQ	0.463	89.5	171	204
400	Soft O2	0.970	30.9	236	313
1600	FlatIP	0.833	359.8	941	1278
1600	IVFPQ	0.330	251.4	669	751
1600	Soft O2	0.917	393.7	985	1172
6400	FlatIP	0.733	4722.8	4458	4969
6400	IVFPQ	0.227	1072.1	2182	3021
6400	Soft O2	0.803	5245.7	4061	4503