

BIMS-LEGAL: Dual-Store Soft Binding for Recovering Prior Legal Advice under Same-Domain Interference

Linrui Xu^a, Linrui Han^{b,c,*}

^a*University of Winnipeg, 515 Portage Avenue, Winnipeg, R3B 2E9, Manitoba, Canada*

^b*Ministry of Education Laboratory of Philosophy and Social Sciences – The CUPL Data Law Lab, China University of Political Science and Law, No. 25 Xitucheng Road, Haidian District, Beijing, 100088, China*

^c*The Institute for Data Law, China University of Political Science and Law, No. 25 Xitucheng Road, Haidian District, Beijing, 100088, China*

Abstract

Recovering prior client-specific advice from a legal consultation archive is difficult: dense retrieval may rank a relevant question highly while omitting the lawyer’s answer, an *episode incompleteness* failure under same-domain interference. BIMS-LEGAL indexes Chinese consultation logs in a dual store—a fast episodic turn index with session identifiers and a slower BIRCH semantic store. Soft O2 binds session siblings by attenuated score inheritance (β s) rather than hard score copy. Exact replay is reported only as a surface-form reference; the main channels are paraphrase, follow-up, and advice-recall. On a shared turn-level index Soft O2 improves over dense FlatIP, with hard hydration and shuffled-*sid* controls separating binding from score copying and membership noise. Gains are largest where a partial episode hit can pull in its answer sibling. Soft O2-C and gated Hybrid extend binding to BIRCH clusters under a Mix protocol with off-session gold; when query and gold already share a session, Soft O2 is preferred. We report Answer Hit, Episode Completeness, graded nDCG with fixed gold-session IDCG, a complete/incomplete/miss failure taxonomy, and paired significance tests.

Keywords: legal information retrieval, consultation memory, soft episode

*Corresponding author. Postal address: No. 25 Xitucheng Road, Haidian District, Beijing 100088, China.

Email addresses: xu-181@webmail.uwinnipeg.ca (Linrui Xu),
linrui_han@163.com (Linrui Han)

1. Introduction

A junior associate may ask a firm assistant: “Last month you advised this client on terminating a fixed-term contract—what remedy did we recommend?” The system returns several near-identical “contract termination” questions, ranks the stored user utterance highly, and never surfaces the partner’s written advice. The subsequent reply may sound fluent yet be grounded in the wrong episode—or in none. In consultation archives, questions and answers share specialized terminology; neighbouring matters are topically adjacent yet legally distinct; and the professionally useful evidence unit is the full consultation *episode*.

Legal informatics has long studied statutory interpretation, case retrieval, and advisory systems Martinez-Gil (2023); Ashley (2017). Recent legal large language models (LLMs) improve statutory and fact-pattern answering Yue et al. (2023); Huang et al. (2023); Yue et al. (2024). Classical legal information retrieval (IR) mainly targets *external* authority—statutes, cases, or community answers—as documents Askari et al. (2024); Louis et al. (2024); Liu et al. (2022); Zhang et al. (2024). Memory-augmented dialogue systems recover evidence from evolving conversation logs Wu et al. (2024); Maharana et al. (2024); Packer et al. (2023). The gap addressed here is *internal consultation memory*: retrieving prior client-specific advice from a growing same-domain archive under lexical interference.

BIMS-LEGAL (Brain-Inspired Memory System for legal consultation archives) separates a fast episodic pathway—timestamped turn embeddings with speaker role and session identifier *sid*—from a slow semantic pathway based on Balanced Iterative Reducing and Clustering using Hierarchies (BIRCH) with deferred consolidation Zhang et al. (1996). Complementary Learning Systems (CLS) theory McClelland et al. (1995); Kumaran et al. (2016) is used only as an architectural analogy for that dual-pathway split. Under same-domain interference, dense turn retrieval in a vector database often returns a question turn while its answer sibling remains outside top- k —a system-level episodic memory failure we term *episode incompleteness*, distinct from a full session miss and exacerbated by IVFPQ score collapse and context fragmentation at moderate archive sizes.

The *algorithmic contributions* are soft binding operators that complete episodes without hard score copy:

- 36 • **Soft O2 (session soft binding).** When a turn is retrieved with score
37 s , siblings sharing sid inherit $\max(s_{t'}, \beta \cdot s)$, surfacing answer turns
38 under question-hit / answer-miss while preserving rank competition
39 among candidates.
- 40 • **Soft O2-C and gated Hybrid (cluster soft binding).** The same
41 soft-inheritance rule applies within BIRCH clusters when gold evidence
42 is not co-session with the query; Hybrid triggers cluster binding only
43 on direct dense hits to avoid displacing same-session siblings.

44 Two **implementation operators** support large shared-store experiments
45 but are not standalone algorithmic contributions:

- 46 • **O1 (exact FlatIP indexing).** At hundreds-to-thousands of turns,
47 IVFPQ product quantization is under-trained and collapses Answer
48 Hit; FlatIP restores faithful turn scores so Soft O2 inherits uncorrupted
49 sibling signals.
- 50 • **O3 (bulk-load consolidation).** Deferred BIRCH and disk writes
51 make archive construction tractable without changing retrieval seman-
52 tics.

53 We ask four research questions:

- 54 • **RQ1.** Under same-domain interference, how prevalent is episode in-
55 completeness (question hit / answer miss) relative to full session miss?
- 56 • **RQ2.** Do Soft O2 and Soft O2-C address episode incompleteness under
57 measured index distortion, and what incremental gains does Soft O2
58 provide over FlatIP and hard parent hydration?
- 59 • **RQ3.** On the *same* LegalEp /LegalMem stores, when is the slow
60 (BIRCH) pathway effective? Does Soft O2-C help under standard same-
61 session gold, and under a fair Mix of same-/cross-session gold?
- 62 • **RQ4.** Across CAIL multi-turn and LegalEp paraphrase /advice-recall
63 channels, how large and session-tied are Soft O2 gains, and when do
64 Soft O2-C /Hybrid help beyond Soft O2?

65 *Contributions..*

- 66 1. A dual-store memory architecture for Chinese legal consultation archives,
67 separating fast episodic turn recovery from slower BIRCH semantic in-
68 tegration under same-domain interference.
- 69 2. Soft O2, a turn-level session soft-binding operator that recovers missing
70 answer siblings without hard parent hydration or session-max score
71 copy.
- 72 3. A three-candidate score-competition account of Soft O2 using the gap
73 ratio $\rho = s_d/s_h$, with default $\beta=0.98$ chosen by validation sweep.
- 74 4. Turn-level lexical and fusion controls (BM25-turn, dense+BM25 RRF,
75 cross-encoder) under the same evidence unit as Soft O2, with joint
76 episode indexing treated as a coarser alternative design rather than a
77 turn-binding baseline.
- 78 5. Soft O2-C and gated Hybrid on Mix cross-session reconstructions, graded
79 nDCG with fixed gold-session IDCG, a failure taxonomy, and paired
80 significance tests.
- 81 6. O1 exact FlatIP and O3 bulk-load consolidation as implementation
82 operators for the shared-store ablations.

83 Section 2 reviews related work. Section 3 presents BIMS-LEGAL opera-
84 tors and the Soft O2 attenuation analysis. Section 4 details corpora, metrics,
85 and protocols. Section 5 reports results. Section 6 discusses implications and
86 limitations. Section 7 concludes.

87 2. Related Work

88 2.1. Legal IR and long-horizon dialogue memory

89 Legal QA and retrieval systems typically ground generation in statutes,
90 cases, or published advice Louis et al. (2024); Askari et al. (2024); Martinez-
91 Gil (2023); Ashley (2017). Conversational legal case retrieval and event-aware
92 similar-case ranking further show how interaction context and structured le-
93 gal knowledge shape IR effectiveness Liu et al. (2022); Zhang et al. (2024).
94 Those pipelines optimize access to external authority. Consultation memory
95 retrieval targets a different evidence source: *prior consultation episodes* inside
96 an agent’s own archive. Confusing two clients’ histories can produce fluent
97 yet professionally inappropriate responses even when statute retrieval is cor-
98 rect. Domain LLMs such as DISC-LawLLM, Lawyer LLaMA, and LawLLM

improve statutory and advisory generation Yue et al. (2023); Huang et al. (2023); Yue et al. (2024), but they still need a reliable memory index when prior advice must be recovered from a growing same-domain log.

Legal consultation archives add a further difficulty: high lexical overlap among topically adjacent matters, with the professionally useful evidence unit being the full question-answer episode. This setting motivates *dual-store* memory architectures that keep fast episodic turn retrieval separate from slower semantic clustering, and session-aware soft binding operators that recover answer turns under same-domain interference without collapsing rank competition among unrelated candidates.

2.2. Dual-store memory and episodic retrieval

Dual-store designs distinguish fast episodic encoding from slower semantic integration McClelland et al. (1995); Kumaran et al. (2016); Tulving et al. (1972); Atkinson and Shiffrin (1971). Memory architectures for dialogue agents extend context via hierarchical paging, summary banks, graphs, and neuro-symbolic retrieval Packer et al. (2023); Zhong et al. (2024); Rasmussen et al. (2025); Gutiérrez et al. (2024); Modarressi et al. (2023); Wang et al. (2024); Lewis et al. (2020); Wu et al. (2025), but most treat retrieved units as flat passages rather than *multi-turn consultation episodes* whose answer turn may rank below an unrelated question turn sharing legal vocabulary. Long-MemEval and LoCoMo evaluate long-horizon recall and QA over personal conversation logs Wu et al. (2024); Maharana et al. (2024), while broader memory-agent surveys emphasize persistence beyond the context window Wu et al. (2025); Tan et al. (2025). Recent conversational RAG work stresses reuse of historical search evidence across turns Mo et al. (2026). BIMS-LEGAL positions its dual-store contribution at this intersection: episodic turn embeddings with *sid* for same-session recovery (Soft O2), and BIRCH centroids with Soft O2-C for cross-session recovery when query and gold evidence occupy different sessions but share thematic structure Zhang et al. (1996). Under dense same-domain interference, a question turn can occupy top- k while its answer sibling is absent. Parent-document hydration and session-max aggregation recover that sibling by hard score copy; Soft O2 instead injects attenuated inherited scores so episode completion does not erase rank competition among unrelated candidates.

2.3. Session-aware retrieval

Parent-document hydration, session aggregation, and joint session indexing are standard session-aware mechanisms for attaching sibling context to a dense hit. Soft episode binding (Soft O2) inherits $\beta \cdot s$ for siblings and keeps ranking competition among candidates. Parent-document and hierarchical retrieval similarly attach child evidence to a parent unit Callan et al. (1994); Dai et al. (2019); Zaheer et al. (2017); Karpukhin et al. (2020); Soft O2 differs by using attenuated turn-level score propagation rather than hard hydration or atomic session documents. Lexical BM25 Robertson and Zaragoza (2009), dense passage retrieval Karpukhin et al. (2020), dense-sparse Reciprocal Rank Fusion Cormack et al. (2009), and cross-encoder reranking Nogueira et al. (2019); Chen et al. (2024) provide complementary controls outside session membership. Exact flat indexes and product-quantized ANN backends Johnson et al. (2019) further affect whether sibling binding is applied over faithful turn scores. Table 1 summarizes the session-structure contrasts used in the main grids.

Table 1: Session-aware mechanisms compared in this work.

Mechanism	Unit	Expansion	Type	Rank
Parent hydration	Session	Insert siblings	Hard	No
Session-max	Session	Session max	Hard	Same*
Joint episode index	Session doc	Atomic	N/A	N/A
Soft O2	Turn+ <i>sid</i>	Soft βs	Soft	Yes
Soft O2-C	Turn+cluster	Soft $\beta_c s$	Soft	Yes

*On the reported corpora, session-max rankings coincide with hard parent hydration under full-sibling expansion; main tables report a single hard-expansion baseline. Soft O2 denotes session soft binding; Soft O2-C denotes cluster soft binding on the BIRCH pathway.

In short, the dual-store split is the system design, and Soft O2 /Soft O2-C are the operators evaluated on Chinese legal consultation archives.

3. Method: BIMS-LEGAL

3.1. Architecture overview

Figure 1 summarizes the pipeline. BIMS-LEGAL maintains complementary episodic and semantic stores over the same consultation archive.




Figure 1: BIMS-LEGAL dual-store pipeline. (A) Episodic turn indexing and BIRCH semantic clustering (O3 bulk-load). (B) Query-time FlatIP retrieval with Soft O2 session binding and Soft O2-C cluster binding; O1 and O3 are implementation operators. (C) Evaluation order: shared-store ablations, Soft O2 grids, and Mix Soft O2-C.

159 **Definition 1** (Episodic and semantic elements). *An episodic element is*
160 $e_i = (\mathbf{v}_i, t_i, c_i, \phi_i, \text{sid}_i)$, *where* $\mathbf{v}_i \in \mathbb{R}^d$ *is an embedding,* t_i *a timestamp,* c_i
161 *conversational context (including speaker role),* ϕ_i *surface features, and* sid_i
162 *the consultation session identifier. A semantic element is* $s_j = (\mathbf{c}_j, C_j, \kappa_j, \tau_j)$
163 *with centroid* $\mathbf{c}_j$, *member set* C_j , *label* κ_j , *and formation statistics* τ_j . Write
164 $\text{Sib}(t) = \{t' : \text{sid}_{t'} = \text{sid}_t\} \setminus \{t\}$ *for session siblings and* $\text{Clus}(t) = \{t' : \text{cid}_{t'} =$
165 $\text{cid}_t\} \setminus \{t\}$ *for BIRCH-cluster siblings.*

166 **Definition 2** (Dual-store retrieval map). *Given query* q *with embedding*
167 $\mathbf{v}_q$, *base episodic retrieval returns a scored candidate set* $\mathcal{R}_0(q) = \{(t, s_t)\}$
168 *under Eq. (1). Soft O2 produces* $\mathcal{R}_{\text{O2}}(q) = \text{SoftBind}(\mathcal{R}_0(q); \text{Sib}, \beta)$ *(Sec-*
169 *tion 3.3). Soft O2-C produces* $\mathcal{R}_{\text{O2C}}(q) = \text{SoftBind}(\mathcal{R}_0(q); \text{Clus}, \beta_c)$ *(Sec-*
170 *tion 3.5). Gated Hybrid composes session binding then cross-session cluster*
171 *binding (Algorithm 2). The returned list is the top-k of the bound set under*
172 *the rewritten scores.*

173 3.2. Base scoring and implementation operators (O1, O3)

174 For query q with embedding $\mathbf{v}_q$, episodic base retrieval scores a memory
175 turn m by

$$\text{score}(m) = \alpha S(\mathbf{v}_q, \mathbf{v}_m) + (1 - \alpha) R(t_m), \quad (1)$$

176 where S is cosine similarity on ℓ_2 -normalized embeddings and

$$R(t_m) = \max\left(0, 1 - \frac{\Delta_t}{30 \text{ days}}\right), \quad (2)$$

177 with Δ_t the elapsed time between query time and the turn timestamp. Unless
178 stated otherwise, $\alpha=0.7$. Eq. (1) is the ranking score used throughout the
179 legal grids; Soft O2 /Soft O2-C rewrite sibling scores after this fusion step
180 and do not replace Eq. (1).

181 *O1: Exact FlatIP indexing..* Let $\hat{S}$ denote the similarity induced by an ANN
182 backend. At hundreds to a few thousand turns, FAISS IVFPQ is often under-
183 trained, so $\|\hat{S} - S\|$ is large enough to reorder turns and break Soft O2
184 inheritance. O1 forces $\hat{S} = S$ via exact FlatIP (inner-product search on
185 ℓ_2 -normalized embeddings). Quantization error depresses Answer Hit (Ta-
186 ble A.3: IVFPQ AH 0.330 at $M=1600$ on LegalEp-DISC exact). O1 is an
187 engineering prerequisite for Soft O2 evaluation, not a standalone retrieval
188 algorithm.

189 *O3: Bulk-load consolidation..* During bulk construction, O3 replaces the per-
190 insert update ($e_i \mapsto \text{BIRCH}(e_i); \text{Flush}$) by a deferred operator $\text{Finalize} =$
191 $\text{BIRCH}(\{e_i\}_{i=1}^N) \circ \text{Flush}$, executed once after all episodic writes. O3 does
192 not change retrieval semantics; it makes $M \geq 400$ shared-store experiments
193 runnable (≈ 11 min wall-clock for $M=400$; amortized ≈ 0.825 s/turn over
194 ~ 800 turns). Quality–latency curves to $M=6400$ appear in Section 5.6.

195 3.3. Soft O2: session soft binding

196 *Motivation..* Under same-domain interference, a retrieved question turn can
197 score highly while its answer sibling remains outside top- k —episode incom-
198 pleteness compounded by IVFPQ score collapse and context fragmentation.
199 Soft O2 completes episodes without hard score copy.

200 **Definition 3** (Soft binding operator). *Given a scored set $\mathcal{R} = \{(t, s_t)\}$, a*
201 *sibling relation $\mathcal{N}(\cdot)$, and attenuation $\beta \in (0, 1]$, define*

$$\text{SoftBind}(\mathcal{R}; \mathcal{N}, \beta) : \quad s_{t'} \leftarrow \max(s_{t'}, \beta \cdot s_t) \quad \text{for all } t \in \mathcal{R}, t' \in \mathcal{N}(t). \quad (3)$$

Algorithm 1 Soft O2 session soft binding

Require: Query q , base hits $\mathcal{R}_0$, top- k , attenuation β

```
1:  $\mathcal{R} \leftarrow \mathcal{R}_0$ 
2: for each hit  $t \in \mathcal{R}_0$  with score  $s_t$  do
3:   for each sibling  $t' \in \text{Sib}(t)$  absent from  $\mathcal{R}$  do
4:     insert  $(t', \beta \cdot s_t)$  into  $\mathcal{R}$ 
5:   end for
6: end for
7:                                      $\triangleright$  Completeness pass after fusion with Eq. (1)
8: for each session  $sid$  represented in  $\mathcal{R}$  do
9:    $s_{\max} \leftarrow \max\{s_u : u \in \mathcal{R}, sid_u = sid\}$ 
10:  for each  $t' \in \text{Sib}$  of  $sid$  still absent do
11:    insert  $(t', \beta \cdot s_{\max})$  into  $\mathcal{R}$ 
12:  end for
13: end for
14: return top- $k$  of  $\mathcal{R}$  by descending score
```

202 *Soft O2 is $\text{SoftBind}(\cdot; \text{Sib}, \beta)$ with default $\beta=0.98$. Hard parent hydration is*
203 *the special case $\beta=1$ with forced insertion of all siblings; session-max replaces*
204 *each member score by $\max_{t \in sid} s_t$.*

205 Algorithm 1 states Soft O2 as sibling expansion followed by a complete-
206 ness pass and top- k truncation, matching the public implementation.

207 3.4. Soft O2 attenuation and score competition

208 The attenuation β controls how strongly a retrieved turn promotes its
209 session siblings. Availability of the injected sibling rises with β , while fidelity
210 of the direct hit falls as $\beta \rightarrow 1$. We characterize that trade-off with a three-
211 candidate score-competition model and choose β by validation sweep rather
212 than by a closed-form optimum. A zero-margin hinge objective on $\beta \in (0, 1]$
213 is uninformative for this choice: its rank term vanishes for all $\beta < 1$.

214 **Definition 4** (Three-candidate score-competition model). *For a query whose*
215 *dense retrieval directly hits a gold turn h with fused score $s_h > 0$, let a be*
216 *the gold answer sibling and let d be the strongest non-session distractor in a*
217 *wide candidate pool, with scores s_a and s_d . Soft O2 rewrites*

$$s'_h = s_h, \quad s'_a = \max(s_a, \beta s_h), \quad s'_d = s_d, \quad (4)$$

Table 2: Gap-ratio statistics on LegalEp $M=3000$ stores (exact-channel query proxy; strongest measured non-session distractor). Default β is chosen by the AH sweep in Table A.4.

Corpus	ρ_{50}	ρ_{95}	Cover@0.90	Default β
LegalEp-DISC	0.746	0.854	0.981	0.98
LegalEp-Lawyer	0.662	0.815	0.998	0.98

and defines the gap ratio $\rho = s_d/s_h$. Because d is the strongest measured distractor, conditions based on ρ are conservative sufficient conditions for beating that distractor, not exact top- k boundary conditions for every competitor.

The binding analysis is most relevant in the soft-injection regime $s_a \leq \beta s_h$, where $s'_a = \beta s_h$.

Proposition 1 (Feasible attenuation band). *Under Definition 4, assume the soft-injection case $s_a \leq \beta s_h$. Then (i) availability against the measured distractor requires $\beta > \rho$; (ii) rank preservation of the direct hit over the injected sibling requires $\beta < 1$, and more generally $s'_a < s_h$; (iii) when $\rho < 1$, any $\beta \in (\rho, 1)$ satisfies both conditions in the active injection case. If $s_a > s_h$ already, rank preservation depends on the observed sibling score rather than on β alone.*

Proof. If $s_a \leq \beta s_h$, then $s'_a = \beta s_h$. Availability follows from $\beta s_h > s_d$, i.e. $\beta > \rho$. Rank preservation requires $s'_a < s_h$, which reduces to $\beta < 1$ when $s_h > 0$. $\square$

Soft-rank visualization and sweep.. On a validation grid we also plot the soft pairwise surrogate

$$\mathcal{L}_{\text{sr}}(\beta; \rho) = -\log \sigma(\tau(\beta - \rho)) - \gamma \log \sigma(\tau(1 - \beta)), \quad (5)$$

with logistic σ , temperature τ , and rank weight γ . The curve is a visualization of the availability–ranking trade-off under the measured ρ pool; it is not used as a fitted optimum. Default $\beta=0.98$ is taken from the AH sweep in Table A.4, where $\{0.90, 0.95, 0.98\}$ form a stable high-AH band on CAIL follow-up and LegalEp paraphrase as well as exact grids. Advice-recall uses the same default without a separate gap-ratio fit.

242 Fair controls include hard parent hydration, session-max aggregation, and
 243 shuffled *sid*. On two-turn LegalEp episodes, hard hydration and session-max
 244 coincide under full-sibling expansion, so tables report a single hard-expansion
 245 baseline. Concatenating a full question–answer episode into one retrieval unit
 246 (BM25-joint or dense joint-episode documents) changes the evidence unit
 247 from turns to documents; we treat that design as an alternative indexing
 248 choice rather than a Soft O2 turn-binding baseline (Section 4).

249 3.5. Semantic consolidation and Soft O2-C

250 Semantic consolidation follows a BIRCH-style dual-phase procedure Zhang
 251 et al. (1996).

252 **Definition 5** (BIRCH dual-phase consolidation). *Let $\{\mathbf{c}_j\}$ be current cen-*
 253 *troids and θ_H, θ_M the assign /merge thresholds. **Phase 1 (assign).** For*
 254 *a new embedding $\mathbf{v}$, set $j^* = \arg \max_j S(\mathbf{v}, \mathbf{c}_j)$; if $S(\mathbf{v}, \mathbf{c}_{j^*}) \geq \theta_H$ then*
 255 *$C_{j^*} \leftarrow C_{j^*} \cup \{\mathbf{v}\}$ and refresh $\mathbf{c}_{j^*}$, else spawn a new cluster. **Phase 2***
 256 *(merge). Periodically, merge pairs with $S(\mathbf{c}_i, \mathbf{c}_j) \geq \theta_M$ and purge near-*
 257 *duplicates above a redundancy threshold. Under O3, both phases run inside*
 258 **Finalize.** *Each turn t receives a cluster id cid_t , inducing $\text{Clus}(t)$ in Defini-*
 259 *tion 1.*

260 **Definition 6** (Soft O2-C). *Soft O2-C is $\text{SoftBind}(\cdot; \text{Clus}, \beta_c)$ with default*
 261 *$\beta_c=0.95 < \beta$, reflecting noisier thematic neighbourhoods than session mem-*
 262 *bership. A per-cluster sibling budget B_c caps $|\text{Clus}(t) \cap \text{Injected}|$ so large*
 263 *thematic clusters cannot flood top- k .*

264 Gated Hybrid (Algorithm 2) injects only *cross-session* cluster siblings
 265 relative to sessions already covered by dense /Soft O2 hits, and triggers
 266 cluster binding only on direct dense hits. This prevents thematic confusers
 267 from displacing same-session siblings already recovered by Soft O2. Soft O2
 268 applies when query and gold share *sid*. Soft O2-C is the cluster-pathway
 269 operator when gold evidence lies in another session of the same thematic
 270 cluster.

271 4. Corpora and Evaluation Protocol

272 Figure 2 sketches the shared-corpus design used throughout the legal eval-
 273 uation. Gold needles and same-domain distractors occupy one store. Each
 274 query probes that shared archive under fixed interference conditions. Pri-
 275 vately rebuilt haystacks are avoided so that system comparisons isolate the
 276 retrieval operator.

Algorithm 2 Gated Hybrid Soft O2 + Soft O2-C

Require: Base hits $\mathcal{R}_0$, budgets (k, B_c) , attenuations (β, β_c)

- 1: $\mathcal{R} \leftarrow \text{SoftBind}(\mathcal{R}_0; \text{Sib}, \beta)$ $\triangleright$ session pathway first
 - 2: $\mathcal{S}_{\text{cov}} \leftarrow \{sid_t : t \in \text{top-}k(\mathcal{R})\}$
 - 3: $\mathcal{T}_{\text{trig}} \leftarrow \{t \in \mathcal{R}_0\}$ $\triangleright$ direct dense hits only
 - 4: **for** each trigger $t \in \mathcal{T}_{\text{trig}}$ **do**
 - 5: $s_{\text{max}} \leftarrow \max\{s_u : u \in \mathcal{R}, cid_u = cid_t\}$
 - 6: **for** each $t' \in \text{Clus}(t)$ with $sid_{t'} \notin \mathcal{S}_{\text{cov}}$, up to budget B_c **do**
 - 7: $s_{t'} \leftarrow \max(s_{t'}, \beta_c \cdot s_{\text{max}})$
 - 8: **end for**
 - 9: **end for**
 - 10: **return** top- k of $\mathcal{R}$
-

277 *4.1. Datasets and query channels*

278 Multi-turn authenticity is evaluated on CAIL2024 consultation dialogues.
279 Gold needles are authentic preliminary and final dialogues, with conversation
280 turns completed by label turns where applicable. Query channels cover the
281 first user turn (U1), a later in-dialogue user turn (Uk-followup), the last user
282 turn (U-last), and constrained paraphrases (U-para). Evidence defaults to
283 assistant turns in the target session.

284 LegalEp-DISC and LegalEp-Lawyer rebuild public DISC-Law-SFT and
285 Lawyer-LLaMA subsets into consultation-episode corpora. Each episode is
286 a natural (question, answer) pair with session identifier *sid*. Reconstruction
287 applies length bounds and legal-content filters, excludes exam prompts, re-
288 moves near-duplicates, separates needles from distractors under a fixed seed,
289 and retains an audit subsample. After filtering, each LegalEp corpus retains
290 3500 passed sessions (500 needles, 2500 distractors at the $M=3000$ primary
291 tier).

292 *Query-channel policy..* Exact replay uses the stored user question and is
293 reported only as a *diagnostic* upper bound on surface-form recovery. Primary
294 claims use paraphrase, Uk-followup /U-last, and advice-recall queries whose
295 surface form is not identical to the stored question text; paraphrase and
296 advice-recall templates are generated under a fixed seed and audited on a
297 subsample. LegalMem-MT reuses CAIL multi-turn gold under the same store
298 for same-session Soft O2 vs Soft O2-C contrasts and for Mix reconstructions.

299 Archive size is varied with gold needles held fixed. Unless stated otherwise,




Figure 2: Shared-corpus protocol. Gold needles and same-domain distractors share one store; queries probe multiple channels with turn-level systems and significance tests. Exact replay is diagnostic; paraphrase, follow-up, and advice-recall are primary channels.

300 Soft O2 transfer tables use the $M \approx 3000$ tier; O1–O2 ablations use $M=400$,
 301 $S=300$.

302 4.2. Metrics and baselines

303 Primary metrics are Answer Hit@ k (AH), Episode Completeness@ k (EC),
 304 and normalized Discounted Cumulative Gain@ k (nDCG) at $k=10$. AH is
 305 binary: whether a gold answer turn appears in the top- k . EC is the mean
 306 coverage of gold turns within the target session. Session Hit@ k is reported
 307 only for diagnosis. For nDCG we use fixed graded relevance over the full
 308 gold session: answer turns $rel=1.0$, other same-session turns $rel=0.5$, all
 309 other turns $rel=0.0$. The ideal denominator (IDCG) is computed from all
 310 gold-relevant turns in the target session and truncated at k , so every system
 311 on a query shares the same IDCG. Answer Hit and Episode Completeness
 312 are the main Soft O2 outcomes; nDCG is reported with the fixed gold-session
 313 IDCG so that systems remain comparable on each query. Hard expansion
 314 can raise nDCG when copied siblings occupy early ranks, so AH and EC are
 315 read together with nDCG.

316 The primary hypothesis family is Soft O2 vs FlatIP on the nine CAIL
 317 and LegalEp transfer channels; for that family we report bootstrap 95% CIs,
 318 McNemar mid- p , and Holm-adjusted p -values (Table A.2). Hard-hydration,
 319 CE, and Soft O2-C /Hybrid Mix contrasts are secondary.

320 Turn-level dense FlatIP (O1) is the base dense retriever. Soft O2 applies
 321 soft sibling score inheritance on the episodic pathway; Soft O2-C applies
 322 the same rule within BIRCH clusters. Hard parent hydration expands a hit
 323 to all siblings by copying the trigger score. Lexical controls include BM25
 324 over turns and dense+BM25 Reciprocal Rank Fusion (RRF) under the same
 325 turn-level evidence unit. BM25-joint and dense joint-episode indexes retrieve
 326 pre-concatenated episode documents and therefore answer a coarser retrieval
 327 problem; they are discussed as alternative designs in Section 6.3 rather than
 328 Soft O2 turn-binding baselines. A cross-encoder (CE) control reranks the
 329 top-30 FlatIP candidates with BAAI/bge-reranker-v2-m3. Shuffled *sid*,
 330 a β sweep, and IVFPQ provide membership, attenuation, and quantization
 331 controls.

332 4.3. Fair Mix protocol for Soft O2-C

333 On fully same-session LegalMem grids Soft O2 already recovers gold an-
 334 swers that share *sid*, leaving Soft O2-C little exclusive headroom and expos-
 335 ing it to thematic confusers. To study off-session recovery we rebuild the *same*
 336 LegalEp and LegalMem archives under a Mix protocol: 70% of gold episodes
 337 are split into query-side S_a and evidence-side S_b with distinct session identi-
 338 fiers, while 30% remain intact. Distractors are unchanged. FlatIP, Soft O2,
 339 Soft O2-C, and gated Hybrid share unrestricted dense retrieval on this store;
 340 we report overall AH and stratum AH on `cross_session` vs `same_session`
 341 queries.

342 *Cluster co-membership under Mix.* After BIRCH consolidation, each cross-
 343 session S_a/S_b pair is assigned a shared cluster identifier. This construction
 344 uses split metadata available only when building the evaluation store; it does
 345 not inject labels at query time. The intent is to test Soft O2-C when thematic
 346 co-membership is present. If recovery still fails, the failure is attributable to
 347 binding rather than accidental cluster separation. Natural same-cluster rates
 348 without this construction are lower (Section 6.3). Soft O2’s session member-
 349 ship rule is unchanged. LegalEp Mix uses advice-recall queries; LegalMem
 350 Mix uses mid-split Uk-followup queries. LongMemEval /LoCoMo numbers
 351 from prior runs on this codebase (QA correctness 70.7% / 68.2%) supply
 352 long-horizon memory context only Wu et al. (2024); Maharana et al. (2024).

353 5. Experiments

354 Unless noted, the embedding model is `nomic-embed-text-v2-moe` Nuss-
 355 baum et al. (2025), $k=10$, and the main $M\approx 3000$ grids use seed 42. A
 356 five-seed check on a smaller shared store ($M=400$) is reported later; robust-
 357 ness across alternative encoders is left open (Section 6.3). Primary contrasts
 358 appear in figures; full numeric grids are in the appendix. Each main table
 359 maps to a released JSON artifact in the accompanying repository.

360 5.1. Implementation operators and Soft O2 ablation

361 Table 3 reports the shared-store ablation on DISC-Law and Lawyer-LLaMA
 362 ($M=400$, $S=300$). IVFPQ baseline Answer Hit is 0.540 / 0.397. O1 alone
 363 yields a small gain on DISC (0.567) and a near-tie on Lawyer (0.397), confirm-
 364 ing that exact indexing is a necessary implementation prerequisite but insuffi-
 365 cient alone. O1+Soft O2 raises Answer Hit to 0.780 / 0.680 (+0.240 / +0.283
 366 over baseline) and Episode Completeness to 0.888 / 0.840. O3 is enabled for
 367 all configurations in this ladder as a construction prerequisite; it does not
 368 alter retrieval scores.

Table 3: O1+Soft O2 ablation on the shared store ($M=400$, $S=300$, $k=10$). EC denotes Episode Completeness (gold-turn coverage); AH denotes Answer Hit. O3 bulk-load is on for all rows. Boldface: best AH and EC within each corpus. Graded nDCG for the $M\approx 3000$ grids appears in Table 6.

Corpus	System	EC@10	AH@10
DISC-Law	IVFPQ baseline	0.770	0.540
	+ O1 FlatIP	0.782	0.567
	+ O1+O2 Soft O2	0.888	0.780
Lawyer-LLaMA	IVFPQ baseline	0.698	0.397
	+ O1 FlatIP	0.698	0.397
	+ O1+O2 Soft O2	0.840	0.680

369 5.2. RQ1 failure taxonomy

370 We classify each query into *complete* (answer evidence in top- k), *incom-*
 371 *plete* (target session partially retrieved but answer missing), or *session miss*
 372 (no target-session turn in top- k). Table 4 contrasts FlatIP, Soft O2, hard
 373 hydration, and shuffled *sid* on primary non-exact channels. Soft O2 converts
 374 incompleteness into complete retrieval while session-miss stays nearly fixed;

Table 4: RQ1 failure taxonomy ($k=10$). Boldface: best complete / lowest incomplete within each corpus-channel block among FlatIP, Soft O2, and shuffled *sid* (hard hydration shown as a score-copy envelope). Soft O2 cuts incompleteness without raising session miss; shuffled *sid* undoes the gain.

Corpus / channel	System	Complete	Incomplete	Session miss
DISC / u-param	FlatIP	70.8%	26.8%	1.6%
	Soft O2	89.1%	8.7%	1.8%
	Hard hydr.	97.2%	0.0%	2.8%
	Shuffled O2	66.4%	31.4%	2.2%
DISC / advice	FlatIP	48.0%	17.0%	35.0%
	Soft O2	60.8%	3.8%	35.4%
	Hard hydr.	65.0%	0.0%	35.0%
	Shuffled O2	45.6%	19.2%	35.2%
Lawyer / u-param	FlatIP	71.5%	25.3%	0.8%
	Soft O2	94.0%	3.8%	1.6%
	Hard hydr.	98.0%	0.0%	2.0%
	Shuffled O2	72.9%	25.5%	1.6%
CAIL / Uk	FlatIP	45.8%	54.0%	0.2%
	Soft O2	92.8%	7.0%	0.2%
	Hard hydr.	98.3%	1.5%	0.2%
	Shuffled O2	34.7%	65.2%	0.2%

shuffled *sid* restores high incompleteness, so the gain is session-tied rather than generic score inflation. Hard hydration can also drive incompleteness near zero by copying trigger scores, but it is not Soft O2’s binding rule and can over-promote siblings (Section 5.4).

5.3. Turn-level Soft O2 controls

Primary Soft O2 claims use a shared turn-level index. Table 5 reports Answer Hit for FlatIP, Soft O2, hard hydration, and shuffled *sid* on the primary Soft O2 transfer channels. Soft O2 beats FlatIP on every non-exact channel; shuffled *sid* collapses toward or below FlatIP; hard hydration is competitive on CAIL AH but Soft O2 remains stronger on LegalEp paraphrase /advice and retains higher EC in the full grids (Tables A.5–A.7). BM25-turn is a same-granularity lexical control (Table A.8): Soft O2 exceeds it on CAIL Uk /U-last (0.698/0.762 vs 0.482/0.413). Joint episode / document-level indexes change evidence granularity and are not Soft O2 protocol competitors.

Table 5: Primary Soft O2 controls (AH@10, $M \approx 3000$). Boldface: best AH within each row among FlatIP / Soft O2 / Hard / Shuffled. Soft O2 is the session-binding claim; shuffled *sid* tests membership dependence; hard hydration is score-copy expansion.

Corpus / channel	FlatIP	Soft O2	Hard	Shuffled
CAIL / U1	0.570	0.788	0.802	0.452
CAIL / Uk	0.270	0.698	0.720	0.230
CAIL / U-last	0.245	0.762	0.790	0.200
LegalEp-DISC / u-para	0.547	0.648	0.592	0.469
LegalEp-DISC / advice	0.342	0.428	0.402	0.310
LegalEp-Lawyer / u-para	0.568	0.759	0.596	0.592
LegalEp-Lawyer / advice	0.332	0.472	0.394	0.348

389 5.4. Soft O2 on session (CAIL and LegalEp)

390 Figure 3 and Table 5 summarize AH@10 on authentic CAIL multi-turn
391 channels at $M \approx 3000$ under O1 FlatIP. Soft O2 lifts Answer Hit over FlatIP
392 on every channel (U1 0.788 vs 0.570**; Uk 0.698 vs 0.270**; U-last 0.762 vs
393 0.245**). Shuffled *sid* removes the gain (U1 0.452; Uk 0.230; U-last 0.200),
394 tying the improvement to correct episode membership (RQ1–RQ2). Hard
395 expansion can raise AH slightly above Soft O2 (e.g., U1 0.802 vs 0.788)
396 while lowering EC (U1 0.578 vs Soft O2 0.664). Soft O2 therefore behaves
397 as a rank-preserving binder: siblings enter under attenuated scores, and
398 episode coverage remains higher than under hard score copy. Table 6 shows
399 the Soft O2>FlatIP ordering under graded nDCG with gold-session IDCG,
400 while incompleteness falls sharply (Table 4). Full numeric grids appear in
401 Tables A.5–A.7.

402 Figure 4 reports LegalEp transfer on the same AH grids. On LegalEp-
403 DISC exact replay Soft O2 and FlatIP stay near-tied (0.638 vs 0.640), where
404 both turns are already recoverable for many queries—consistent with treat-
405 ing exact as diagnostic. On Lawyer exact Soft O2 improves AH to 0.752
406 from FlatIP 0.692 (McNemar mid- $p=0.024^*$). Paraphrase yields the clearest
407 LegalEp gains: Soft O2 reaches 0.648 vs FlatIP 0.547 on DISC and 0.759 vs
408 0.568 on Lawyer, exceeding hard expansion on both corpora (mid- $p<0.001^{**}$
409 on both corpora). Advice-recall queries omit the stored question surface form:
410 Soft O2 raises AH to 0.428 from FlatIP 0.342 on DISC and to 0.472 from
411 0.332 on Lawyer, again above hard expansion (mid- $p<0.001^{**}$), while shuffled
412 *sid* falls toward or below FlatIP. Graded nDCG mirrors the Soft O2>FlatIP
413 ordering on paraphrase and advice-recall (Table 6). Detailed grids are in

Table 6: Graded nDCG@10 with gold-session IDCG (answer $rel=1.0$, other same-session turns $rel=0.5$) and incompleteness rates on primary non-exact channels. Boldface: Soft O2 better than FlatIP within each channel on nDCG. Soft O2 also exceeds shuffled *sid*.

Corpus	Channel	System	nDCG@10	Incomplete
CAIL	Uk	FlatIP	0.445	54.0%
CAIL	Uk	Soft O2	0.793	7.0%
CAIL	Uk	Shuffled O2	0.369	65.2%
CAIL	U-last	FlatIP	0.420	57.3%
CAIL	U-last	Soft O2	0.791	5.3%
CAIL	U-last	Shuffled O2	0.348	71.2%
LegalEp-DISC	u-para	FlatIP	0.681	26.8%
LegalEp-DISC	u-para	Soft O2	0.759	8.7%
LegalEp-DISC	u-para	Shuffled O2	0.627	31.4%
LegalEp-DISC	advice	FlatIP	0.466	17.0%
LegalEp-DISC	advice	Soft O2	0.524	3.8%
LegalEp-DISC	advice	Shuffled O2	0.436	19.2%
LegalEp-Lawyer	u-para	FlatIP	0.744	25.3%
LegalEp-Lawyer	u-para	Soft O2	0.844	3.8%
LegalEp-Lawyer	u-para	Shuffled O2	0.705	25.5%
LegalEp-Lawyer	advice	FlatIP	0.483	23.4%
LegalEp-Lawyer	advice	Soft O2	0.582	2.2%
LegalEp-Lawyer	advice	Shuffled O2	0.454	25.0%

414 Tables A.6–A.7.

415 Cross-encoder reranking is a strong neural control at the same turn-
416 level evidence unit: FlatIP+CE reranks the top-30 FlatIP candidates with
417 **bge-reranker-v2-m3**. Table 7 shows Soft O2 ahead of FlatIP+CE on CAIL
418 multi-turn channels and on LegalEp-Lawyer exact, while FlatIP+CE leads on
419 LegalEp-DISC exact—consistent with Soft O2’s largest gains arising where a
420 partial episode hit can bind its answer sibling rather than where surface-form
421 overlap already recovers both turns.

422 5.5. Soft O2-C and gated Hybrid under Mix

423 When gold answers already share *sid* with the query, Soft O2 is sufficient
424 and cluster binding is unnecessary. Table 8 asks the off-session question
425 under the Mix protocol of Section 4.3: 70% cross-session Sa/Sb gold and
426 30% intact same-session gold, with FlatIP, Soft O2, Soft O2-C, and gated
427 Hybrid sharing unrestricted dense retrieval.

CAIL2024 multi-turn results ($M \approx 3000$, turn-level systems)




Figure 3: CAIL2024 multi-turn results at tier $M \approx 3000$. Soft O2 improves AH and EC over FlatIP; shuffled *sid* removes the gain; hard expansion can maximize AH at a modest EC cost.

Table 7: Cross-encoder control (AH@10). FlatIP+CE reranks the top-30 FlatIP candidates with **bge-reranker-v2-m3**. Boldface: best AH within each row.

Corpus / channel	FlatIP	FlatIP+CE	Soft O2
CAIL / U1	0.585	0.688	0.788
CAIL / Uk	0.265	0.375	0.698
CAIL / U-last	0.228	0.342	0.762
LegalEp-DISC / exact	0.572	0.722	0.638
LegalEp-Lawyer / exact	0.636	0.708	0.752

On the cross-session stratum Soft O2-C improves over Soft O2 on LegalEp-Lawyer (0.491 vs 0.446) and LegalMem-MT (0.537 vs 0.503): cluster binding can surface split-episode gold that session membership alone misses. Gated Hybrid converts that mechanism into an overall gain on LegalEp-Lawyer (0.534 vs Soft O2 0.496; mid- $p=0.010^*$) and a directional gain on LegalMem-MT (0.646 vs 0.618), while preserving most same-session hits (0.560 / 0.867). On LegalEp-DISC the Hybrid edge is negligible (0.486 vs 0.482). Ungated Soft O2-C can raise cross-session hits yet displace intact same-session evidence, which is why Hybrid triggers cluster binding only on direct dense hits. Same-session Soft O2 vs Soft O2-C contrasts remain Soft O2-led (Appendix Appendix A, Table A.1).

McNemar mid- p on paired Answer Hit; $^*p<0.05$, $^{**}p<0.01$. Cluster co-membership of Sa/Sb pairs follows the Mix protocol in Section 4.3.

LegalEp ($M = 3000$): Soft O2 gains are strongest on paraphrase and remain positive on advice-recall




Figure 4: LegalEp AH@10 ($M=3000$, 500 needles). Soft O2 gains are strongest on paraphrase and remain positive on advice-recall; exact transfer is corpus-dependent and treated as diagnostic.




Figure 5: Scale curve on DISC-Law exact. Soft O2 keeps an AH edge over FlatIP as M grows; warm-query latency scales with archive size; IVFPQ quality collapses at moderate M .

5.6. Scale behaviour

Figure 5 and Table A.3 report quality and warm-query latency on LegalEp-DISC exact ($Q=100$, three repeats; $M \in \{100, 400, 1600, 6400\}$). Soft O2 retains an AH advantage at each tier; IVFPQ collapses at $M=1600$ (AH 0.330), motivating O1; absolute latency grows with M . We treat O1 FlatIP as an engineering prerequisite for Soft O2 at these scales and avoid claiming deployment readiness beyond the measured curve.

End-to-end QA audit (summary). A dual-protocol QA audit ($N=270$ per corpus; generator **qwen3:14b**, independent judge **qwen3:32b**) reports judged answerability under Soft O2 retrieval alone; it is not a paired FlatIP vs

Table 8: Mix evaluation of Soft O2-C and gated Hybrid (same store, unrestricted dense; AH@10). Gold is 70% cross-session Sa/Sb + 30% intact same-session. Boldface: best overall AH and best cross-session AH within each corpus. Significance vs Soft O2: * $p < 0.05$, ** $p < 0.01$.

Corpus	System	AH	AH _{cross}	AH _{same}	mid- p vs Soft O2
LegalEp-DISC Mix	FlatIP	0.478	0.477	0.480	—
	Soft O2	0.482	0.443	0.573	—
	Soft O2-C	0.410	0.443	0.333	0.0004**
	Hybrid	0.486	0.469	0.527	0.8099
LegalEp-Lawyer Mix	FlatIP	0.462	0.454	0.480	—
	Soft O2	0.496	0.446	0.613	—
	Soft O2-C	0.464	0.491	0.400	0.0857
	Hybrid	0.534*	0.523	0.560	0.0105*
LegalMem-MT Mix	FlatIP	0.534	0.526	0.553	—
	Soft O2	0.618	0.503	0.887	—
	Soft O2-C	0.500	0.537	0.413	0.0000**
	Hybrid	0.646	0.551	0.867	0.1837

451 Soft O2 generation experiment. Full protocol and stratified results appear in
 452 Section 6.2. On a smaller shared store ($M=400$, five seeds), Soft O2 AH on
 453 LegalEp-DISC is 0.851 ± 0.008 versus FlatIP 0.698 ± 0.040 .

454 6. Discussion

455 6.1. Findings and scope

456 RQ1–RQ2 show that episode incompleteness is a substantial FlatIP fail-
 457 ure mode under same-domain interference, and that Soft O2 converts many
 458 incomplete retrievals into complete ones without shrinking session-miss rates
 459 (Table 4). O1 restores score fidelity; Soft O2 supplies the main gain through
 460 soft sibling inheritance; O3 makes large shared stores tractable. The gap-
 461 ratio analysis (Section 3.4) motivates keeping β close to but below unity; the
 462 AH sweep peaks in $\{0.9, 0.95, 0.98\}$ rather than at hard hydration ($\beta=1$). On
 463 CAIL, Soft O2 raises Answer Hit by large margins (U1 +0.218, Uk +0.428,
 464 U-last +0.517) with collapsing shuffled-*sid* controls, and keeps higher EC
 465 than hard expansion despite occasional AH ties. After Holm adjustment on
 466 the Soft O2 vs FlatIP family, the non-exact CAIL and LegalEp gains re-
 467 main significant, while LegalEp exact replay is mixed (Table A.2). LegalEp

paraphrase and advice-recall, rather than exact replay, support the main transfer result. Graded nDCG agrees with the Soft O2>FlatIP ordering (Table 6). Under Mix, gated Hybrid recovers additional cross-session gold on LegalEp-Lawyer with a directional lift on LegalMem-MT (Table 8); Soft O2 remains preferable when gold already shares *sid* (Table A.1). Turn-level BM25, RRF, and CE place Soft O2 among standard IR alternatives Robertson and Zaragoza (2009); Cormack et al. (2009); Nogueira et al. (2019); joint-episode indexes are a coarser design choice outside this protocol.

6.2. RAG application perspective

Consultation-memory retrieval is judged downstream by whether an LLM can answer from recovered evidence rather than confabulate. We audit this end-to-end with a dual-protocol pipeline on Soft O2 retrieval only: generator `qwen3:14b` and independent judge `qwen3:32b` are distinct models. Per corpus we pool $N=270$ judged instances as P1 $n=120$ plus P2 $n=150$ (folder `legal_qa_n270`). Table 9 summarizes pooled Evidence-Answerability Rate (EAR) and stratified independent-judge scores. In the complete-episode stratum—both question and answer turns in top- k —judged answerability is higher in this single-system audit (pooled EAR 0.867 on DISC and 0.859 on Lawyer; Wilson 95% CI $[0.83, 0.90]$ / $[0.82, 0.89]$). Follow-up queries remain the hardest stratum (independent-judge EAR 0.517 / 0.480), consistent with episodic incompleteness persisting even when surface-form overlap is low.

A blind two-annotator legal audit ($n=60$ /corpus) yields Cohen’s $\kappa=0.71$ / 0.68 and separates retrieval from generation failures (retrieval-failure share 0.283 / 0.317). Wrong-session contamination rates stay low (0.083 / 0.100), suggesting that complete episode retrieval under Soft O2 is associated with lower rates of a common RAG failure mode in which the model answers from a neighbouring client’s matter. Hallucinated-legal-basis rates (0.117 / 0.133) remain non-zero, indicating that grounding in retrieved episodes does not guarantee statutory correctness—a scope boundary for deployment. AH and EAR measure evidence availability and judged answerability respectively; statutory correctness, citation validity, and professional liability lie outside this evaluation.

6.3. Limitations

LegalEp episodes are consultation pairs; multi-turn authenticity remains reserved for CAIL. The Mix protocol assigns Sa/Sb pairs a shared cluster id

Table 9: End-to-end QA audit ($N=270$ per corpus; generator `qwen3:14b`, independent judge `qwen3:32b`). EAR: Evidence-Answerability Rate (judge score ≥ 0.7). Stratified judge scores: $n=90$ /corpus across exact, paraphrase, and follow-up protocols.

Corpus	Pooled EAR	Stratified judge	Human κ
DISC-Law	0.867	0.757	0.71
Lawyer-LLaMA	0.859	0.729	0.68

so Soft O2-C can be tested when thematic co-membership is present (Section 4.3); unsupervised BIRCH without that construction would not receive split-pair metadata at build time. Soft O2 depends on correct session identifiers; Soft O2-C depends on cluster quality. Joint-episode document indexes change the evidence unit from turns to concatenated episodes; comparing them under a matched online-update and latency protocol is left to future work. The ρ statistics use an exact-channel query proxy and the strongest measured distractor; paraphrase and advice-recall gap laws may shift the feasible band slightly, though the AH sweep keeps a stable high- β region on those channels. All main grids use one multilingual embedding model and a fixed seed at $M \approx 3000$; five-seed variance is reported only for the smaller $M=400$ store. LLM paraphrases need human audit. Statutory correctness, citation validity, and professional liability remain unevaluated. Because many channels are tested, Holm-adjusted p -values accompany the Soft O2 vs FlatIP family (Table A.2); Hybrid Mix gains outside LegalEp-Lawyer are treated as directional.

Latency and deployment scope. Scale curves cover $M \in \{100, 400, 1600, 6400\}$ under a fixed hardware stack (Table A.3). At $M=6400$, warm-query p50 latency exceeds 4s (4061 ms for Soft O2; 4458 ms for FlatIP)—unacceptable for millisecond interactive search. BIMS-LEGAL should therefore be positioned as **high-precision offline or quasi-real-time consultation memory recovery**—e.g., pre-meeting case review, audit replay, or batch RAG over a firm’s advice archive—not as a sub-millisecond search engine. Future work will combine approximate nearest-neighbour indexes (e.g., HNSW) with Soft O2 and Soft O2-C binding layered on top of faithful candidate retrieval, preserving episode completion while reducing query latency. Behavior at roughly 10^5 sessions is not measured, and we do not claim asymptotic scalability from O3 wall-clock alone.

532 7. Conclusion

533 BIMS-LEGAL recovers prior advice from Chinese legal consultation
534 archives under same-domain interference with a dual store: an episodic turn
535 index and a BIRCH semantic store. Soft O2 is the main algorithmic con-
536 tribution: attenuated session binding completes episodes without hard score
537 copy. Soft O2-C and gated Hybrid extend the same idea to thematic clus-
538 ters when Mix reconstructions place some gold off-session. Attenuation β
539 is chosen by validation sweep under a three-candidate score-competition ac-
540 count; O1 FlatIP and O3 bulk-load are the implementation operators that
541 make the shared-store ablations feasible. On CAIL multi-turn and LegalEp
542 paraphrase /advice-recall channels, Soft O2 improves Answer Hit over FlatIP,
543 keeps higher episode completeness than hard expansion on CAIL, and raises
544 graded nDCG while cutting incompleteness; shuffled-*sid* controls collapse
545 the gains. Under Mix, gated Hybrid recovers additional cross-session gold on
546 LegalEp-Lawyer while Soft O2 remains preferable on intact same-session gold.
547 An end-to-end Soft O2 QA audit associates complete episode retrieval with
548 judged answerability and low wrong-session contamination, without claim-
549 ing paired generation superiority over FlatIP. Turn-level BM25, RRF, and
550 CE place Soft O2 among standard IR alternatives; joint-episode indexing
551 remains a coarser design outside this protocol.

552 Declarations

553 *Funding*

554 Computational resources were provided by the Data Law Laboratory,
555 China University of Political Science and Law.

556 *Conflict of interest*

557 The authors declare no competing interests.

558 *Ethics approval*

559 This study uses publicly released research corpora and excludes confiden-
560 tial client data. The system is a research prototype for consultation-memory
561 retrieval.

562 *Data availability*

563 Code, processed evaluation corpora, and primary result JSON files are
564 available at <https://github.com/Kilimajaro/soft-episode-binding-legal-memory> (archived release to accompany camera-ready). Upstream
565 corpora remain under their original licenses: CAIL2024 consultation tracks;
566 Hugging Face ShengbinYue/DISC-Law-SFT and Skepsun/lawyer_llama_data. Processed LegalEp /LegalMem artifacts are derived research releases;
567
568 redistributors should consult each upstream licence before reuse.
569

570 *Code availability*

571 The BIMS-LEGAL implementation (O1 FlatIP, Soft O2, O3 bulk-load, Soft O2-C), corpus builders, Mix constructors, and table/figure scripts ship
572 in the same repository, with one JSON artifact per main table.
573

574 *Author contributions*

575 L.X. (University of Winnipeg) designed BIMS-LEGAL and the Soft O2 /Soft O2-C evaluation, implemented the system, ran the experiments, and wrote the manuscript. L.H. (corresponding author; CUPL Data Law Lab and Institute for Data Law) provided methodology guidance, revision, and supervision.
576
577
578
579

580 **References**

- 581 Ashley, K.D., 2017. Artificial Intelligence and Legal Analytics: New Tools
582 for Law Practice in the Digital Age. Cambridge University Press.
- 583 Askari, A., Aliannejadi, M., Abolghasemi, A., Kanoulas, E., Verberne, S.,
584 2024. Answer retrieval in legal community question answering, in: Advances in Information Retrieval (ECIR), Springer. pp. 477–485.
- 585
- 586 Atkinson, R.C., Shiffrin, R.M., 1971. The control of short-term memory.
587 Scientific American 225, 82–91.
- 588 Callan, J.P., Lu, Z., Xu, S., 1994. Passage-level evidence in document retrieval, in: Proceedings of the 17th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval, pp. 302–310.
589
590
591

- 592 Chen, J., Xiao, S., Zhang, P., Luo, K., Lian, D., Liu, Z., 2024. M3-
593 embedding: Multi-linguality, multi-functionality, multi-granularity text
594 embeddings through self-knowledge distillation, in: Findings of the associ-
595 ation for computational linguistics: ACL 2024, pp. 2318–2335.
- 596 Cormack, G.V., Clarke, C.L.A., Buettcher, S., 2009. Reciprocal rank fusion
597 outperforms condorcet and individual rank learning methods, in: Proceed-
598 ings of the 32nd International ACM SIGIR Conference on Research and
599 Development in Information Retrieval, pp. 758–759.
- 600 Dai, Z., Xiong, C., Liu, Z., Callan, J., 2019. Context-aware sentence/passage
601 encoding for answering factoid questions, in: Proceedings of the 2019 Con-
602 ference on Empirical Methods in Natural Language Processing and the
603 9th International Joint Conference on Natural Language Processing, pp.
604 241–251.
- 605 Gutiérrez, B.J., Shu, Y., Gu, Y., Yasunaga, M., Su, Y., 2024. Hipporag:
606 Neurobiologically inspired long-term memory for large language models.
607 Advances in neural information processing systems 37, 59532–59569.
- 608 Huang, Q., Tao, M., Zhang, C., An, Z., Jiang, C., Chen, Z., Wu, Z., Feng,
609 Y., 2023. Lawyer llama technical report. [arXiv:2305.15062](https://arxiv.org/abs/2305.15062).
- 610 Johnson, J., Douze, M., Jégou, H., 2019. Billion-scale similarity search with
611 GPUs. IEEE Transactions on Big Data 7, 535–547.
- 612 Karpukhin, V., Oguz, B., Min, S., Lewis, P., Wu, L., Edunov, S., Chen,
613 D., Yih, W.t., 2020. Dense passage retrieval for open-domain question
614 answering, in: Proceedings of the 2020 Conference on Empirical Methods
615 in Natural Language Processing (EMNLP), pp. 6769–6781.
- 616 Kumaran, D., Hassabis, D., McClelland, J.L., 2016. What learning systems
617 do intelligent agents need? complementary learning systems theory up-
618 dated. Trends in cognitive sciences 20, 512–534.
- 619 Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küt-
620 tler, H., Lewis, M., Yih, W.t., Rocktäschel, T., et al., 2020. Retrieval-
621 augmented generation for knowledge-intensive nlp tasks. Advances in neu-
622 ral information processing systems 33, 9459–9474.

- 623 Liu, B., Wu, Y., Zhang, F., Liu, Y., Wang, Z., Li, C., Zhang, M., Ma, S., 2022.
624 Query generation and buffer mechanism: Towards a better conversational
625 agent for legal case retrieval. *Information Processing & Management* 59,
626 103051.
- 627 Louis, A., van Dijck, G., Spanakis, G., 2024. Interpretable long-form legal
628 question answering with retrieval-augmented large language models, in:
629 Proceedings of the AAAI Conference on Artificial Intelligence, pp. 22266–
630 22275.
- 631 Maharana, A., Lee, D.H., Tulyakov, S., Bansal, M., Barbieri, F., Fang, Y.,
632 2024. Evaluating very long-term conversational memory of llm agents, in:
633 Proceedings of the 62nd Annual Meeting of the Association for Computa-
634 tional Linguistics (Volume 1: Long Papers), pp. 13851–13870.
- 635 Martinez-Gil, J., 2023. A survey on legal question–answering systems. *Com-
636 puter Science Review* 48, 100552.
- 637 McClelland, J.L., McNaughton, B.L., O’Reilly, R.C., 1995. Why there are
638 complementary learning systems in the hippocampus and neocortex: in-
639 sights from the successes and failures of connectionist models of learning
640 and memory. *Psychological review* 102, 419.
- 641 Mo, F., Gao, Y., Wu, Z., Liu, X., Chen, P., Li, Z., Wang, Z., Li, X., Jiang,
642 M., Nie, J.Y., 2026. Leveraging historical information to boost retrieval-
643 augmented generation in conversations. *Information Processing & Man-
644 agement* 63, 104449.
- 645 Modarressi, A., Imani, A., Fayyaz, M., Schütze, H., 2023. Ret-llm: Towards
646 a general read-write memory for large language models. *arXiv preprint
647 arXiv:2305.14322* .
- 648 Nogueira, R., Yang, W., Cho, K., Lin, J., 2019. Multi-stage document rank-
649 ing with BERT. *arXiv:1910.14497*.
- 650 Nussbaum, Z., et al., 2025. Nomic embed multimodal: Interleaved text, im-
651 age, and datasets for general-purpose vision-language embeddings. Model
652 card: `nomic-embed-text-v2-moe` (Nomic AI). Embedding model used for
653 turn encoding in this work.

- 654 Packer, C., Wooders, S., Lin, K., Fang, V., Patil, S.G., Stoica, I., Gonzalez,
655 J.E., 2023. Memgpt: Towards llms as operating systems. arXiv preprint
656 arXiv:2310.08560 .
- 657 Rasmussen, P., Paliychuk, P., Beauvais, T., Ryan, J., Chalef, D., 2025. Zep:
658 A temporal knowledge graph architecture for agent memory. arXiv preprint
659 arXiv:2501.13956 .
- 660 Robertson, S., Zaragoza, H., 2009. The probabilistic relevance framework:
661 BM25 and beyond. Foundations and Trends in Information Retrieval 3,
662 333–389.
- 663 Tan, H., Zhang, Z., Ma, C., Chen, X., Dai, Q., Dong, Z., 2025. Mem-
664 bench: Towards more comprehensive evaluation on the memory of llm-
665 based agents, in: Findings of the Association for Computational Linguis-
666 tics: ACL 2025, pp. 19336–19352.
- 667 Tulving, E., et al., 1972. Episodic and semantic memory. Organization of
668 memory 1, 1.
- 669 Wang, Y., Gao, Y., Chen, X., Jiang, H., Li, S., Yang, J., Yin, Q., Li, Z.,
670 Li, X., Yin, B., et al., 2024. Memoryllm: Towards self-updatable large
671 language models. arXiv preprint arXiv:2402.04624 .
- 672 Wu, D., Wang, H., Yu, W., Zhang, Y., Chang, K.W., Yu, D., 2024. Long-
673 memeval: Benchmarking chat assistants on long-term interactive memory.
674 arXiv preprint arXiv:2410.10813 .
- 675 Wu, Y., Liang, S., Zhang, C., Wang, Y., Zhang, Y., Guo, H., Tang, R.,
676 Liu, Y., 2025. From human memory to ai memory: A survey on memory
677 mechanisms in the era of llms. arXiv:2504.15965.
- 678 Yue, S., Chen, W., Wang, S., Li, B., Shen, C., Liu, S., Zhou, Y., Xiao, Y.,
679 Yun, S., Huang, X., Wei, Z., 2023. Disc-lawllm: Fine-tuning large language
680 models for intelligent legal services. arXiv:2309.11325.
- 681 Yue, S., Liu, S., Zhou, Y., Shen, C., Wang, S., Xiao, Y., Li, B., Yun, S.,
682 Shen, X., Chen, W., et al., 2024. Lawllm: Intelligent legal system with
683 legal reasoning and verifiable retrieval, in: International Conference on
684 Database Systems for Advanced Applications, Springer. pp. 304–321.

- 685 Zaheer, M., Kottur, S., Ravanbakhsh, S., Poczos, B., Salakhutdinov, R.,
686 Smola, A.J., 2017. Deep sets, in: Advances in Neural Information Process-
687 ing Systems.
- 688 Zhang, T., Ramakrishnan, R., Livny, M., 1996. BIRCH: an efficient data
689 clustering method for very large databases, in: Proceedings of the 1996
690 ACM SIGMOD International Conference on Management of Data, pp. 103–
691 114.
- 692 Zhang, Y., Zhai, S., Meng, Y., Bi, S., Chen, Y., Qi, G., 2024. Event is
693 more valuable than you think: Improving the similar legal case retrieval
694 via event knowledge. *Information Processing & Management* 61, 103729.
- 695 Zhong, W., Guo, L., Gao, Q., Ye, H., Wang, Y., 2024. Memorybank: En-
696 hancing large language models with long-term memory, in: Proceedings of
697 the AAAI Conference on Artificial Intelligence, pp. 19724–19731.

698 **Appendix A. Numeric grids and controls**

699 Same-session Soft O2-C contrasts, significance tests, scale curves, and full
700 Soft O2 numeric grids appear below.

Table A.1: Same-session Soft O2 vs ungated Soft O2-C (AH@10). Soft O2 is preferable when gold already shares *sid* with the query; ungated Soft O2-C can surface thematic confusers. Boldface: better AH within each row. Mix Soft O2-C results appear in Table 8.

Corpus / channel	Soft O2	Soft O2-C
LegalEp-DISC / exact	0.836	0.744
LegalEp-DISC / u-para	0.807	0.730
LegalEp-DISC / advice	0.548	0.490
LegalEp-Lawyer / exact	0.818	0.710
LegalEp-Lawyer / u-para	0.817	0.733
LegalEp-Lawyer / advice	0.546	0.448
LegalMem-MT / U1	0.903	0.712
LegalMem-MT / Uk	0.793	0.385

Table A.2: McNemar mid- p on paired Answer Hit (Soft O2 vs FlatIP / hard hydration / FlatIP+CE). Holm-adjusted p is reported for the nine primary Soft O2 vs FlatIP contrasts.

Corpus / channel	Contrast	ΔAH	mid- p	Holm p
CAIL / Uk	O2 vs FlatIP	+0.428	8.99e-48	8.09e-47
CAIL / Uk	O2 vs Hard hydr.	-0.022	0.243	—
CAIL / U1	O2 vs FlatIP	+0.218	5.19e-18	4.67e-17
CAIL / U1	O2 vs Hard hydr.	-0.013	0.435	—
CAIL / U-last	O2 vs FlatIP	+0.517	5.51e-65	4.96e-64
CAIL / U-last	O2 vs Hard hydr.	-0.028	0.0982	—
CAIL / Uk	O2 vs FlatIP+CE	+0.323	2.50e-26	—
CAIL / U1	O2 vs FlatIP+CE	+0.100	3.51e-05	—
CAIL / U-last	O2 vs FlatIP+CE	+0.420	3.92e-41	—
LegalEp-DISC / exact	O2 vs FlatIP	-0.002	0.934	0.934
LegalEp-DISC / exact	O2 vs Hard hydr.	+0.024	0.302	—
LegalEp-Lawyer / exact	O2 vs FlatIP	+0.060	0.0239	0.191
LegalEp-Lawyer / exact	O2 vs Hard hydr.	+0.140	3.50e-10	—
LegalEp-DISC / u-para	O2 vs FlatIP	+0.101	7.95e-05	6.36e-04
LegalEp-DISC / u-para	O2 vs Hard hydr.	+0.056	0.035	—
LegalEp-Lawyer / u-para	O2 vs FlatIP	+0.191	2.10e-15	1.68e-14
LegalEp-Lawyer / u-para	O2 vs Hard hydr.	+0.163	7.57e-09	—
LegalEp-DISC / advice	O2 vs FlatIP	+0.086	1.25e-04	8.75e-04
LegalEp-DISC / advice	O2 vs Hard hydr.	+0.026	0.215	—
LegalEp-Lawyer / advice	O2 vs FlatIP	+0.140	1.18e-10	9.44e-10
LegalEp-Lawyer / advice	O2 vs Hard hydr.	+0.078	5.72e-04	—

Table A.3: Scale curve on DISC-Law (mean over 3 repeats). Latency is warm-query p50/p95. Boldface: best AH@10 within each M block; lowest p50/p95 within each M block.

M	Mode	AH@10	Build (s)	p50 (ms)	p95 (ms)
100	FlatIP	0.963	3.0	67	80
100	IVFPQ	0.613	25.2	40	54
100	Soft O2	0.970	3.2	63	69
400	FlatIP	0.890	26.5	248	305
400	IVFPQ	0.463	89.5	171	204
400	Soft O2	0.970	30.9	236	313
1600	FlatIP	0.833	359.8	941	1278
1600	IVFPQ	0.330	251.4	669	751
1600	Soft O2	0.917	393.7	985	1172
6400	FlatIP	0.733	4722.8	4458	4969
6400	IVFPQ	0.227	1072.1	2182	3021
6400	Soft O2	0.803	5245.7	4061	4503

Table A.4: Soft O2 AH@10 versus β on the primary shared-corpus evaluation. Peak performance lies in $\{0.9, 0.95, 0.98\}$; aggressive attenuation ($\beta=0.5$) weakens binding.

Corpus	0.5	0.7	0.9	0.95	0.98	1.0
CAIL / Uk	0.323	0.537	0.697	0.700	0.698	0.662
LegalEp-Lawyer / exact	0.608	0.654	0.754	0.754	0.752	0.752
LegalEp-Lawyer / u-para	0.631	0.637	0.717	0.745	0.759	0.749
LegalEp-DISC / u-para	0.551	0.551	0.618	0.648	0.648	0.642

Table A.5: CAIL2024 shared-corpus results ($M \approx 3000$). Boldface marks the best AH@10 and best EC@10 within each channel (ties bolded jointly). **: Soft O2 vs FlatIP, McNemar mid- $p < 0.01$. nDCG@10 uses gold-session IDCg.

Channel	System	AH@10	EC@10	nDCG@10	95% CI (AH)
U1	FlatIP	0.570	0.413	0.554	[0.530,0.610]
	Soft O2	0.788	0.664	0.776	[0.757,0.820]
	Hard hydr.	0.802	0.578	0.864	[0.768,0.835]
	Shuffled O2	0.452	0.363	0.411	[0.410,0.492]
Uk-followup	FlatIP	0.270	0.347	0.445	[0.237,0.307]
	Soft O2	0.698	0.654	0.793	[0.660,0.735]
	Hard hydr.	0.720	0.585	0.847	[0.685,0.753]
	Shuffled O2	0.230	0.320	0.369	[0.198,0.262]
U-last	FlatIP	0.245	0.335	0.420	[0.212,0.280]
	Soft O2	0.762	0.659	0.791	[0.725,0.797]
	Hard hydr.	0.790	0.586	0.854	[0.757,0.822]
	Shuffled O2	0.200	0.315	0.348	[0.172,0.232]

701

Soft O2 leads EC; hard hydration can lead AH.

Table A.6: LegalEp-DISC ($M=3000$, 500 needles). Boldface: best AH@10 and best EC@10 within each channel. * $p < 0.05$, ** $p < 0.01$ (Soft O2 vs FlatIP, McNemar mid- p). Hard hydration coincides with session-max aggregation; only hard hydration is shown. nDCG@10 uses gold-session IDCg.

Channel	System	AH@10	EC@10	nDCG@10
exact	FlatIP	0.640	0.741	0.730
	Soft O2	0.638	0.780	0.813
	Hard hydr.	0.614	0.598	0.886
	Shuffled O2	0.448	0.645	0.666
advice-recall	FlatIP	0.342	0.438	0.466
	Soft O2	0.428	0.503	0.524
	Hard hydr.	0.402	0.405	0.572
	Shuffled O2	0.310	0.421	0.436
u-param	FlatIP	0.547	0.681	0.681
	Soft O2	0.648	0.735	0.759
	Hard hydr.	0.592	0.571	0.827
	Shuffled O2	0.469	0.610	0.627

Table A.7: LegalEp-Lawyer ($M=3000$, 500 needles). Boldface: best AH@10 and best EC@10 within each channel. * $p<0.05$, ** $p<0.01$ (Soft O2 vs FlatIP). Hard hydration coincides with session-max aggregation; only hard hydration is shown. nDCG@10 uses gold-session IDCG.

Channel	System	AH@10	EC@10	nDCG@10
exact	FlatIP	0.692	0.806	0.753
	Soft O2	0.752	0.848	0.869
	Hard hydr.	0.612	0.615	0.897
	Shuffled O2	0.570	0.747	0.705
advice-recall	FlatIP	0.332	0.464	0.483
	Soft O2	0.472	0.540	0.582
	Hard hydr.	0.394	0.403	0.601
	Shuffled O2	0.348	0.467	0.454
u-param	FlatIP	0.568	0.718	0.744
	Soft O2	0.759	0.819	0.844
	Hard hydr.	0.596	0.594	0.864
	Shuffled O2	0.592	0.712	0.705

Table A.8: Turn-level BM25 and dense+BM25 RRF (AH@10). Boldface: higher AH within each row among reported turn-level cells. BM25-joint / joint-episode document indexes are omitted: they change evidence granularity and were not run as Soft O2 protocol competitors.

Corpus	Channel	AH@10
LegalEp-DISC	exact (BM25-turn)	0.802
LegalEp-DISC	advice-recall (BM25-turn)	0.382
LegalEp-DISC	exact (dense+BM25 RRF)	0.780
LegalEp-DISC	advice (dense+BM25 RRF)	0.394
CAIL	U1 / Uk / U-last (BM25-turn)	0.748 / 0.482 / 0.413
LegalEp-Lawyer	exact / advice (BM25-turn)	0.784 / 0.460